

## 12\_feat-importances

June 24, 2022

# CPSC 330

# Applied Machine Learning

## 1 Lecture 12: Feature importances

UBC 2022 Summer

Instructor: Mehrdad Oveisi

### 1.1 Imports

```
[1]: import os
import string
import sys
from collections import deque

import matplotlib.pyplot as plt
import numpy as np
import pandas as pd

sys.path.append("code/.")

import seaborn as sns
from plotting_functions import *
from sklearn import datasets
from sklearn.compose import ColumnTransformer, make_column_transformer
from sklearn.dummy import DummyClassifier, DummyRegressor
from sklearn.ensemble import RandomForestClassifier, RandomForestRegressor
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LogisticRegression, Ridge
from sklearn.model_selection import (
    GridSearchCV,
    RandomizedSearchCV,
    cross_val_score,
```

```

        cross_validate,
        train_test_split,
    )
    from sklearn.pipeline import Pipeline, make_pipeline
    from sklearn.preprocessing import OneHotEncoder, OrdinalEncoder, StandardScaler
    from sklearn.svm import SVC, SVR
    from sklearn.tree import DecisionTreeClassifier
    from utils import *

    %matplotlib inline

```

## 1.2 Learning outcomes

From this lecture, students are expected to be able to:

- Interpret the coefficients of linear regression for ordinal, one-hot encoded categorical, and scaled numeric features.
- Explain why interpretability is important in ML.
- Use `feature_importances_` attribute of `sklearn` models and interpret its output.
- Use `eli5` to get feature importances of non `sklearn` models and interpret its output.
- Apply SHAP to assess feature importances and interpret model predictions.
- Explain force plot, summary plot, and dependence plot produced with shapely values.

## 1.3 Data

In this lecture, we'll be using [Kaggle House Prices dataset](#), the dataset we used in lecture 2. As usual, to run this notebook you'll need to download the data. Unzip the data into a subdirectory called `data`. For this dataset, train and test have already been separated. We'll be working with the train portion in this lecture.

```

[2]: df = pd.read_csv("data/housing-kaggle/train.csv")
    train_df, test_df = train_test_split(df, test_size=0.10, random_state=123)
    train_df.head()

```

```

[2]:      Id  MSSubClass MSZoning  LotFrontage  LotArea Street Alley LotShape  \
302   303           20        RL         118.0   13704   Pave   NaN    IR1
767   768           50        RL          75.0   12508   Pave   NaN    IR1
429   430           20        RL         130.0   11457   Pave   NaN    IR1
1139  1140           30        RL          98.0    8731   Pave   NaN    IR1
558   559           60        RL          57.0   21872   Pave   NaN    IR2

      LandContour Utilities  ... PoolArea PoolQC Fence MiscFeature MiscVal  \
302          Lvl1   AllPub  ...         0    NaN   NaN          NaN         0
767          Lvl1   AllPub  ...         0    NaN   NaN          Shed       1300
429          Lvl1   AllPub  ...         0    NaN   NaN          NaN         0
1139         Lvl1   AllPub  ...         0    NaN   NaN          NaN         0
558          HLS   AllPub  ...         0    NaN   NaN          NaN         0

```

	MoSold	YrSold	SaleType	SaleCondition	SalePrice
302	1	2006	WD	Normal	205000
767	7	2008	WD	Normal	160000
429	3	2009	WD	Normal	175000
1139	5	2007	WD	Normal	144000
558	8	2008	WD	Normal	175000

[5 rows x 81 columns]

- The prediction task is predicting `SalePrice` given features related to properties.
- Note that the **target is numeric**, not categorical.

```
[3]: train_df.shape
```

```
[3]: (1314, 81)
```

### 1.3.1 Let's separate X and y

```
[4]: X_train = train_df.drop(columns=["SalePrice"])
     y_train = train_df["SalePrice"]

     X_test = test_df.drop(columns=["SalePrice"])
     y_test = test_df["SalePrice"]
```

### 1.3.2 Let's identify feature types

```
[5]: drop_features = ["Id"]
     numeric_features = [
         "BedroomAbvGr",
         "KitchenAbvGr",
         "LotFrontage",
         "LotArea",
         "OverallQual",
         "OverallCond",
         "YearBuilt",
         "YearRemodAdd",
         "MasVnrArea",
         "BsmtFinSF1",
         "BsmtFinSF2",
         "BsmtUnfSF",
         "TotalBsmtSF",
         "1stFlrSF",
         "2ndFlrSF",
         "LowQualFinSF",
         "GrLivArea",
         "BsmtFullBath",
         "BsmtHalfBath",
```

```

    "FullBath",
    "HalfBath",
    "TotRmsAbvGrd",
    "Fireplaces",
    "GarageYrBlt",
    "GarageCars",
    "GarageArea",
    "WoodDeckSF",
    "OpenPorchSF",
    "EnclosedPorch",
    "3SsnPorch",
    "ScreenPorch",
    "PoolArea",
    "MiscVal",
    "YrSold",
]

```

```

[6]: ordinal_features_reg = [
    "ExterQual",
    "ExterCond",
    "BsmtQual",
    "BsmtCond",
    "HeatingQC",
    "KitchenQual",
    "FireplaceQu",
    "GarageQual",
    "GarageCond",
    "PoolQC",
]
ordering = [
    "Po",
    "Fa",
    "TA",
    "Gd",
    "Ex",
]
# if N/A it will just impute something, per below
ordering_ordinal_reg = [ordering] * len(ordinal_features_reg)
ordering_ordinal_reg

```

```

[6]: [['Po', 'Fa', 'TA', 'Gd', 'Ex'],
      ['Po', 'Fa', 'TA', 'Gd', 'Ex'],
      ['Po', 'Fa', 'TA', 'Gd', 'Ex'],
      ['Po', 'Fa', 'TA', 'Gd', 'Ex'],
      ['Po', 'Fa', 'TA', 'Gd', 'Ex'],
      ['Po', 'Fa', 'TA', 'Gd', 'Ex'],
      ['Po', 'Fa', 'TA', 'Gd', 'Ex'],
      ['Po', 'Fa', 'TA', 'Gd', 'Ex'],

```

```
['Po', 'Fa', 'TA', 'Gd', 'Ex'],
['Po', 'Fa', 'TA', 'Gd', 'Ex']]
```

```
[7]: ordinal_features_oth = [
    "BsmtExposure",
    "BsmtFinType1",
    "BsmtFinType2",
    "Functional",
    "Fence",
]
ordering_ordinal_oth = [
    ["NA", "No", "Mn", "Av", "Gd"],
    ["NA", "Unf", "LwQ", "Rec", "BLQ", "ALQ", "GLQ"],
    ["NA", "Unf", "LwQ", "Rec", "BLQ", "ALQ", "GLQ"],
    ["Sal", "Sev", "Maj2", "Maj1", "Mod", "Min2", "Min1", "Typ"],
    ["NA", "MnWw", "GdWo", "MnPrv", "GdPrv"],
]
```

```
[8]: categorical_features = list(
    set(X_train.columns)
    - set(numeric_features)
    - set(ordinal_features_reg)
    - set(ordinal_features_oth)
    - set(drop_features)
)
categorical_features
```

```
[8]: ['MasVnrType',
    'Exterior2nd',
    'LandSlope',
    'RoofStyle',
    'Neighborhood',
    'SaleType',
    'Condition2',
    'Exterior1st',
    'SaleCondition',
    'MSZoning',
    'BldgType',
    'RoofMatl',
    'HouseStyle',
    'MiscFeature',
    'MoSold',
    'LotConfig',
    'MSSubClass',
    'Utilities',
    'CentralAir',
    'LotShape',
```

```

'Street',
'Foundation',
'Electrical',
'GarageFinish',
'Condition1',
'Alley',
'LandContour',
'Heating',
'GarageType',
'PavedDrive']

```

```
[9]: from sklearn.compose import ColumnTransformer, make_column_transformer
```

```

numeric_transformer = make_pipeline(SimpleImputer(strategy="median"),
    ↳ StandardScaler())
ordinal_transformer_reg = make_pipeline(
    SimpleImputer(strategy="most_frequent"),
    OrdinalEncoder(categories=ordering_ordinal_reg),
)

ordinal_transformer_oth = make_pipeline(
    SimpleImputer(strategy="most_frequent"),
    OrdinalEncoder(categories=ordering_ordinal_oth),
)

categorical_transformer = make_pipeline(
    SimpleImputer(strategy="constant", fill_value="missing"),
    OneHotEncoder(handle_unknown="ignore", sparse=False),
)

preprocessor = make_column_transformer(
    ("drop", drop_features),
    (numeric_transformer, numeric_features),
    (ordinal_transformer_reg, ordinal_features_reg),
    (ordinal_transformer_oth, ordinal_features_oth),
    (categorical_transformer, categorical_features),
)

```

```
[10]: preprocessor.fit(X_train)
preprocessor.named_transformers_
```

```
[10]: {'drop': 'drop',
'pipeline-1': Pipeline(steps=[('simpleimputer',
SimpleImputer(strategy='median')),
('standardscaler', StandardScaler())]),
'pipeline-2': Pipeline(steps=[('simpleimputer',
SimpleImputer(strategy='most_frequent'))],

```



1139	0.154795	-0.222647	1.358264	-0.171468	-0.773017
558	0.154795	-0.222647	-0.597924	1.289541	0.663680
...	...	...	...	...	...
1041	1.372763	-0.222647	-0.025381	-0.127107	-0.054669
1122	0.154795	-0.222647	-0.025381	-0.149788	-1.491366
1346	0.154795	-0.222647	-0.025381	1.168244	0.663680
1406	-1.063173	-0.222647	0.022331	-0.203265	-0.773017
1389	0.154795	-0.222647	-0.454788	-0.475099	-0.054669

	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea	BsmtFinSF1	...	\
302	-0.512408	0.993969	0.840492	0.269972	-0.961498	...	
767	1.285467	-1.026793	0.016525	-0.573129	0.476092	...	
429	-0.512408	0.563314	0.161931	-0.573129	1.227559	...	
1139	-0.512408	-1.689338	-1.679877	-0.573129	0.443419	...	
558	-0.512408	0.828332	0.598149	-0.573129	0.354114	...	
...	...	...	...	...	...	...	
1041	2.184405	-0.165485	0.743555	0.843281	-0.090231	...	
1122	-2.310284	-0.496757	-1.389065	-0.573129	-0.961498	...	
1346	1.285467	-0.099230	0.888961	-0.573129	-0.314582	...	
1406	1.285467	0.033279	1.082835	-0.573129	0.467379	...	
1389	0.386530	-0.993666	-1.679877	-0.573129	-0.144686	...	

	GarageType_2Types	GarageType_Attchd	GarageType_Basment	\
302	0.0	1.0	0.0	
767	0.0	1.0	0.0	
429	0.0	1.0	0.0	
1139	0.0	0.0	0.0	
558	0.0	1.0	0.0	
...	...	...	...	
1041	0.0	1.0	0.0	
1122	0.0	0.0	1.0	
1346	0.0	1.0	0.0	
1406	0.0	0.0	0.0	
1389	0.0	0.0	0.0	

	GarageType_BuiltIn	GarageType_CarPort	GarageType_Detchd	\
302	0.0	0.0	0.0	
767	0.0	0.0	0.0	
429	0.0	0.0	0.0	
1139	0.0	0.0	1.0	
558	0.0	0.0	0.0	
...	...	...	...	
1041	0.0	0.0	0.0	
1122	0.0	0.0	0.0	
1346	0.0	0.0	0.0	
1406	0.0	0.0	1.0	
1389	0.0	0.0	1.0	



	GarageType_missing	PavedDrive_N	PavedDrive_P	PavedDrive_Y
302	0.0	0.0	0.0	1.0
767	0.0	0.0	0.0	1.0
429	0.0	0.0	0.0	1.0
1139	0.0	0.0	0.0	1.0
558	0.0	0.0	0.0	1.0
...	...	...	...	...
1041	0.0	0.0	0.0	1.0
1122	0.0	0.0	0.0	1.0
1346	0.0	0.0	0.0	1.0
1406	0.0	0.0	0.0	1.0
1389	0.0	0.0	0.0	1.0

[1314 rows x 263 columns]

```
[13]: X_train_enc.shape
```


```
[13]: (1314, 263)
```

### 1.3.3 Feature correlations

- Let's look at the correlations between various features with other features and the target in our encoded data (first row/column).
- In simple terms here is how you can interpret correlations between two variables  $X$  and  $Y$ :
  - If  $Y$  goes up when  $X$  goes up, we say  $X$  and  $Y$  are **positively correlated**.
  - If  $Y$  goes down when  $X$  goes up, we say  $X$  and  $Y$  are **negatively correlated**.
  - If  $Y$  is unchanged when  $X$  changes, we say  $X$  and  $Y$  are **uncorrelated**.

```
[14]: # Get the pairwise correlations between the first 15 columns (including y_train)
cor = pd.concat((y_train, X_train_enc), axis=1).iloc[:, :15].corr()
plt.figure(figsize=(12, 12))
sns.set(font_scale=1)
sns.heatmap(cor, annot=True, cmap=plt.cm.Blues);
```





- Looking at this diagram also tells us the relationship between features.
  - For example, `1stFlrSF` and `TotalBsmtSF` are highly correlated.
  - Do we need both of them?
  - If our model says `1stFlrSF` is very important and `TotalBsmtSF` is very unimportant, do we trust those values?
  - Maybe `TotalBsmtSF` only “becomes important” if `1stFlrSF` is removed.
  - Sometimes the opposite happens: a feature only becomes important if another feature is *added*.

#### 1.4 Feature importance in linear models

- Like logistic regression, with linear regression we can look at the *coefficients* for each feature.
- Overall idea: predicted price = intercept +  $\sum_i \text{coefficient } i \times \text{feature } i$

```
[16]: lr = make_pipeline(preprocessor, Ridge())
      lr.fit(X_train, y_train);
```

Let's look at the coefficients.

```
[17]: lr_coefs = pd.DataFrame(data=lr[1].coef_, index=new_columns,
    ↪columns=["Coefficient"])
lr_coefs.head(20)
```

```
[17]:
```

	Coefficient
BedroomAbvGr	-3723.741570
KitchenAbvGr	-4580.204576
LotFrontage	-1578.664421
LotArea	5109.356718
OverallQual	12487.561839
OverallCond	4855.535334
YearBuilt	4226.684842
YearRemodAdd	324.664715
MasVnrArea	5251.325210
BsmtFinSF1	3667.172851
BsmtFinSF2	583.114880
BsmtUnfSF	-1266.614671
TotalBsmtSF	2751.084018
1stFlrSF	6736.788904
2ndFlrSF	13409.901084
LowQualFinSF	-448.424132
GrLivArea	15988.182407
BsmtFullBath	2299.227266
BsmtHalfBath	500.169112
FullBath	2831.811467

#### 1.4.1 Interpreting coefficients of different types of features.

##### 1.4.2 Ordinal features

- The ordinal features are easiest to interpret.

```
[18]: print(ordinal_features_reg)
```

```
['ExterQual', 'ExterCond', 'BsmtQual', 'BsmtCond', 'HeatingQC', 'KitchenQual',
'FireplaceQu', 'GarageQual', 'GarageCond', 'PoolQC']
```

```
[19]: lr_coefs.loc["ExterQual", "Coefficient"]
```

```
[19]: 4195.671512467826
```

- Increasing by one category of exterior quality (e.g. good -> excellent) increases the predicted price by ~ \$4195.
  - Wow, that's a lot!
  - Remember this is just what the model has learned. It doesn't tell us how the world works.

```
[20]: one_example = X_test[:1]
```

```
[21]: one_example["ExterQual"]
```

```
[21]: 147    Gd
      Name: ExterQual, dtype: object
```

Let's perturb the example and change ExterQual to Ex.

```
[22]: one_example_perturbed = one_example.copy()
      one_example_perturbed["ExterQual"] = "Ex"  # Change Gd to Ex
```

```
[23]: one_example_perturbed
```

```
[23]:      Id  MSSubClass MSZoning  LotFrontage  LotArea Street Alley LotShape  \
147  148          60      RL           NaN     9505   Pave   NaN    IR1

      LandContour Utilities  ... ScreenPorch PoolArea PoolQC Fence MiscFeature  \
147          Lvl   AllPub  ...           0         0   NaN   NaN         NaN

      MiscVal MoSold  YrSold  SaleType  SaleCondition
147         0      5    2010        WD          Normal

[1 rows x 80 columns]
```

```
[24]: one_example_perturbed["ExterQual"]
```

```
[24]: 147    Ex
      Name: ExterQual, dtype: object
```

How does the prediction change after changing ExterQual from Gd to Ex?

```
[25]: print("Prediction on the original example: ", lr.predict(one_example))
      print("Prediction on the perturbed example: ", lr.
      ↪predict(one_example_perturbed))
      print(
          "After changing ExterQual from Gd to Ex increased the prediction by: ",
          lr.predict(one_example_perturbed) - lr.predict(one_example),
      )
```

```
Prediction on the original example: [224795.63596802]
Prediction on the perturbed example: [228991.30748049]
After changing ExterQual from Gd to Ex increased the prediction by:
[4195.67151247]
```

That's exactly the learned coefficient for ExterQual!

```
[26]: lr_coefs.loc["ExterQual", "Coefficient"]
```

```
[26]: 4195.671512467826
```

So our interpretation is correct! - Increasing by one category of exterior quality (e.g. good -> excellent) increases the predicted price by ~ \$4195.

### 1.4.3 Categorical features

- What about the categorical features?
- We have created a number of columns for each category with OHE and each category gets it's own coefficient.

```
[27]: print(categorical_features)
```

```
['MasVnrType', 'Exterior2nd', 'LandSlope', 'RoofStyle', 'Neighborhood',  
'SaleType', 'Condition2', 'Exterior1st', 'SaleCondition', 'MSZoning',  
'BldgType', 'RoofMat1', 'HouseStyle', 'MiscFeature', 'MoSold', 'LotConfig',  
'MSSubClass', 'Utilities', 'CentralAir', 'LotShape', 'Street', 'Foundation',  
'Electrical', 'GarageFinish', 'Condition1', 'Alley', 'LandContour', 'Heating',  
'GarageType', 'PavedDrive']
```

```
[28]: lr_coefs_landslope = lr_coefs[lr_coefs.index.str.startswith("LandSlope")]  
lr_coefs_landslope
```

```
[28]:
```

	Coefficient
LandSlope_Gtl	457.197456
LandSlope_Mod	7420.208381
LandSlope_Sev	-7877.405837

- We can talk about switching from one of these categories to another by picking a “reference” category:

```
[29]: lr_coefs_landslope.loc["LandSlope_Gtl"]
```

```
[29]: Coefficient    457.197456  
Name: LandSlope_Gtl, dtype: float64
```

```
[30]: lr_coefs_landslope - lr_coefs_landslope.loc["LandSlope_Gtl"]
```

```
[30]:
```

	Coefficient
LandSlope_Gtl	0.000000
LandSlope_Mod	6963.010925
LandSlope_Sev	-8334.603292

- If you change the category from LandSlope\_Gtl to LandSlope\_Mod the prediction price goes up by ~ \$6963
- If you change the category from LandSlope\_Gtl to LandSlope\_Sev the prediction price goes down by ~ \$8334

Note that this might not make sense in the real world but this is what our model decided to learn given this small amount of data.

```
[31]: lr_coefs.sort_values(by="Coefficient")
```

```
[31]:
```

	Coefficient
RoofMatl_ClyTile	-191129.774314
Condition2_PosN	-105552.840565
Heating_OthW	-27260.681308
MSZoning_C (all)	-21990.746193
Exterior1st_ImStucc	-19393.964621
...	...
PoolQC	34217.656047
RoofMatl_CompShg	36525.980874
Neighborhood_NridgHt	37532.643270
Neighborhood_StoneBr	39993.978324
RoofMatl_WdShngl	83646.711008

[263 rows x 1 columns]

- For example, the above coefficient says that “If the roof is made of clay tile, the predicted price is \ \$191K less”?
- Do we believe these interpretations??
  - Do we believe this is how the predictions are being **computed**? Yes.
  - Do we believe that this is how the **world works**? No.

```
[32]: # We can see all RoofMatl one hot columns:
lr_coefs[lr_coefs.index.str.startswith("RoofMatl")]
```

```
[32]:
```

	Coefficient
RoofMatl_ClyTile	-191129.774314
RoofMatl_CompShg	36525.980874
RoofMatl_Membran	24537.788381
RoofMatl_Metal	16788.514414
RoofMatl_Roll	8868.963092
RoofMatl_Tar&Grv	7477.664157
RoofMatl_WdShake	13284.152389
RoofMatl_WdShngl	83646.711008

*Note* If you did `drop='first'` (we didn't) then you already have a reference class, and all the values are with respect to that one. The interpretation depends on whether we did `drop='first'`, hence the hassle.

#### 1.4.4 Interpreting coefficients of numeric features

Let's look at coefficients of `PoolArea` and `LotFrontage`.

```
[33]: lr_coefs.loc[["PoolArea", "LotArea", "LotFrontage"]]
```

```
[33]:
```

	Coefficient
PoolArea	2822.370476
LotArea	5109.356718
LotFrontage	-1578.664421

Intuition:

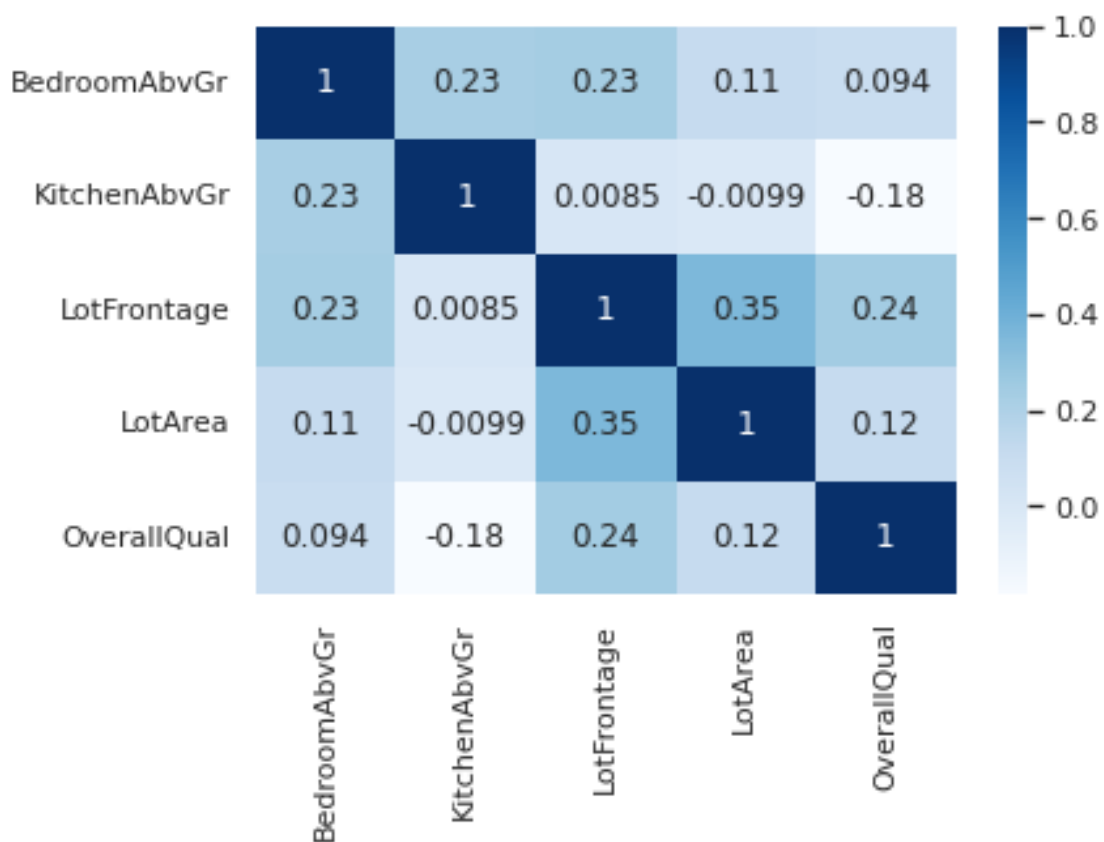
- **Tricky** because numeric features are **scaled**!
- **Increasing** PoolArea by **1 scaled unit** **increases** the predicted price by  $\sim \$2822$ .
- **Increasing** LotFrontage by **1 scaled unit** **decreases** the predicted price by  $\sim \$1578$ .

Does that sound reasonable?

- For PoolArea, yes.
- For LotFrontage, that's surprising. Something positive would have made more sense?

It's not the case here but maybe the problem is that LotFrontage and LotArea are very correlated. LotArea has a larger positive coefficient.

```
[34]: cor = X_train_enc[numeric_features[:5]].corr()
sns.heatmap(cor, annot=True, cmap=plt.cm.Blues);
```



First, let's make sure the predictions behave as expected:

```
[35]: one_example = X_test[:1]
```

```
[36]: one_example
```



```
[36]:      Id  MSSubClass MSZoning  LotFrontage  LotArea Street Alley LotShape  \
147  148           60      RL           NaN     9505   Pave   NaN     IR1

      LandContour Utilities  ... ScreenPorch PoolArea PoolQC Fence MiscFeature  \
147           Lvl   AllPub  ...           0         0   NaN   NaN         NaN

      MiscVal MoSold  YrSold  SaleType  SaleCondition
147         0      5    2010        WD         Normal

[1 rows x 80 columns]
```

Let's **perturb** the example and add 1 to the LotArea.

```
[37]: one_example_perturbed = one_example.copy()
one_example_perturbed["LotArea"] += 1  # add 1 to the LotArea
```

```
[38]: one_example_perturbed
```

```
[38]:      Id  MSSubClass MSZoning  LotFrontage  LotArea Street Alley LotShape  \
147  148           60      RL           NaN     9506   Pave   NaN     IR1

      LandContour Utilities  ... ScreenPorch PoolArea PoolQC Fence MiscFeature  \
147           Lvl   AllPub  ...           0         0   NaN   NaN         NaN

      MiscVal MoSold  YrSold  SaleType  SaleCondition
147         0      5    2010        WD         Normal

[1 rows x 80 columns]
```

Prediction on the original example.

```
[39]: lr.predict(one_example)
```

```
[39]: array([224795.63596802])
```

Prediction on the perturbed example.

```
[40]: lr.predict(one_example_perturbed)
```

```
[40]: array([224796.2040233])
```

- What's the difference between prediction?
- Does the difference make sense given the coefficient of the feature?

```
[41]: lr.predict(one_example_perturbed) - lr.predict(one_example)
```

```
[41]: array([0.56805528])
```

```
[42]: lr_coefs.loc[["LotArea"]]
```

```
[42]:          Coefficient
      LotArea  5109.356718
```

- Why did the prediction only go up by  $\$0.57$  instead of  $\$5109$ ?
- This is an issue of **units**.
  - LotArea is in sqft, but the coefficient is **not** \$5109/sqft **because we scaled the features**.

#### 1.4.5 Example showing how to interpret coefficients of scaled features

- The scaler subtracted the mean and divided by the standard deviation.
- The division actually changed the scale!
- For the unit conversion, we don't care about the subtraction, but only the scaling.

```
[43]: scaler = preprocessor.named_transformers_["pipeline-1"]["standardscaler"]
```

```
[44]: scaler.scale_
```

```
[44]: array([8.21039683e-01, 2.18760172e-01, 2.09591390e+01, 8.99447103e+03,
          1.39208177e+00, 1.11242416e+00, 3.01866337e+01, 2.06318985e+01,
          1.77914527e+02, 4.59101890e+02, 1.63890010e+02, 4.42869860e+02,
          4.42817167e+02, 3.92172897e+02, 4.35820743e+02, 4.69800920e+01,
          5.29468070e+02, 5.18276015e-01, 2.33809970e-01, 5.49298599e-01,
          5.02279069e-01, 1.62604030e+00, 6.34398801e-01, 2.40531598e+01,
          7.40269201e-01, 2.10560601e+02, 1.25388753e+02, 6.57325181e+01,
          6.07432962e+01, 3.03088902e+01, 5.38336322e+01, 4.23249944e+01,
          5.22084645e+02, 1.33231649e+00])
```

```
[45]: lr_scales = pd.DataFrame(
        data=scaler.scale_, index=numeric_features, columns=["Scale"]
    )
    lr_scales.head()
```

```
[45]:          Scale
      BedroomAbvGr    0.821040
      KitchenAbvGr    0.218760
      LotFrontage    20.959139
      LotArea        8994.471032
      OverallQual     1.392082
```

- It seems like LotArea was divided by 8994.471032 sqft.

```
[46]: lr_coefs.loc["LotArea", "Coefficient"]
```

```
[46]: 5109.356718094088
```

- The coefficient tells us that if we increase the **scaled** LotArea by one unit the price would go up by  $\approx \$5109$ .
- One scaled unit represents  $\sim 8994$  sq feet (`lr_scales.loc["LotArea", "Scale"]`)

- So if I increase original `LotArea` by one square foot then the predicted price would go up by this amount:

```
[47]: lr_coefs.loc["LotArea", "Coefficient"] / lr_scales.loc["LotArea", "Scale"]
```

```
[47]: 0.5680552752646643
```

```
[48]: 5109.356718094072 / 8994.471032
```


```
[48]: 0.5680552752814816
```

- This makes much more sense. Now we get the number we got before.
- That said don't read too much into these coefficients without statistical training.

#### 1.4.6 Interim summary

- **Correlation** among features might make coefficients completely **uninterpretable**.
- Fairly **straightforward** to interpret coefficients of **ordinal** features.
- In **categorical** features, it's often helpful to consider **one category as a reference** point and think about relative importance.
- For **numeric** features, relative importance is meaningful **after scaling**.
  - You have to be careful about the scale of the feature when interpreting the coefficients.
- Remember that explaining the model  $\neq$  explaining the data.
  - the **coefficients** tell us only about the **model** and they might **not** accurately reflect the **data**.

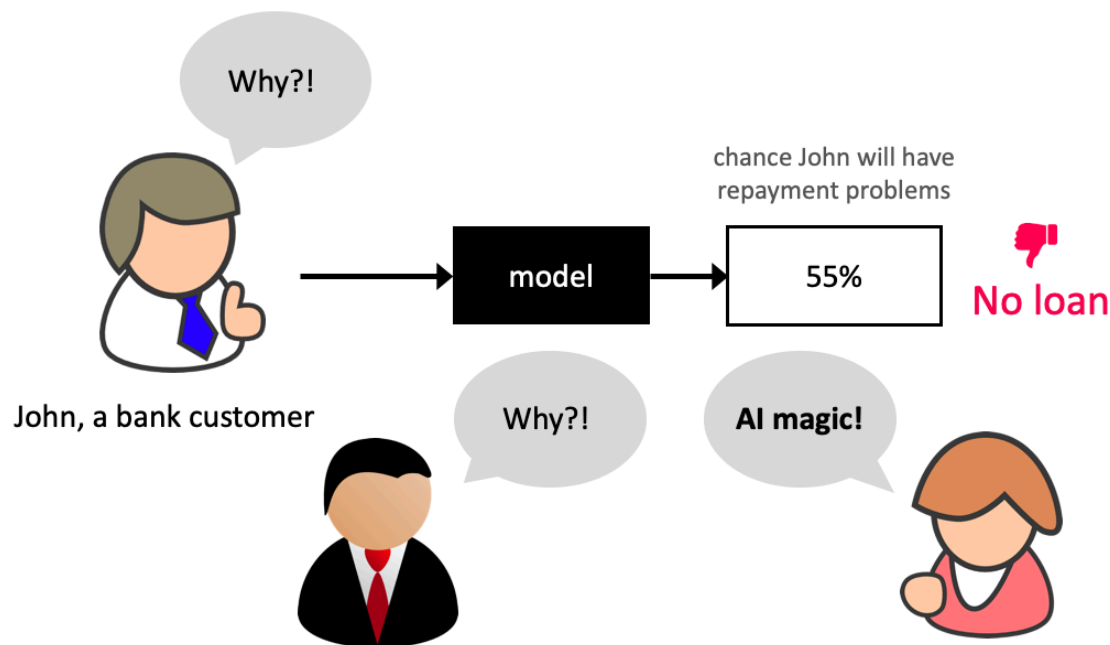
## 1.5 Break (5 min)



## 1.6 Interpretability of ML models: Motivations

### 1.6.1 Why model interpretability?

- Ability to interpret ML models is crucial in many applications such as banking, healthcare, and criminal justice.
- It can be leveraged by domain experts to diagnose systematic errors and underlying biases of complex ML systems.



[Source](#)

### 1.6.2 What is model interpretability?

- In this course, our definition of model interpretability will be looking at **feature importances**.
- There is more to interpretability than feature importances, but it's a good start!
- Resource:
  - [Interpretable Machine Learning](#)
  - [Yann LeCun, Kilian Weinberger, Patrice Simard, and Rich Caruana: Panel debate on interpretability](#)

### 1.6.3 Data

- Let's work with [the adult census data set](#) from last lecture and hw3.

```
[49]: adult_df_large = pd.read_csv("data/adult.csv")
train_df, test_df = train_test_split(adult_df_large, test_size=0.2,
    random_state=42)
train_df_nan = train_df.replace("?", np.NaN)
test_df_nan = test_df.replace("?", np.NaN)
train_df_nan.head()
```

```
[49]:
```

	age	workclass	fnlwgt	education	education.num	\
5514	26	Private	256263	HS-grad	9	
19777	24	Private	170277	HS-grad	9	
10781	36	Private	75826	Bachelors	13	
32240	22	State-gov	24395	Some-college	10	
9876	31	Local-gov	356689	Bachelors	13	

	marital.status	occupation	relationship	race	sex	\
--	----------------	------------	--------------	------	-----	---

5514	Never-married	Craft-repair	Not-in-family	White	Male
19777	Never-married	Other-service	Not-in-family	White	Female
10781	Divorced	Adm-clerical	Unmarried	White	Female
32240	Married-civ-spouse	Adm-clerical	Wife	White	Female
9876	Married-civ-spouse	Prof-specialty	Husband	White	Male

	capital.gain	capital.loss	hours.per.week	native.country	income
5514	0	0	25	United-States	<=50K
19777	0	0	35	United-States	<=50K
10781	0	0	40	United-States	<=50K
32240	0	0	20	United-States	<=50K
9876	0	0	40	United-States	<=50K

```
[50]: numeric_features = ["age", "fnlwgt", "capital.gain", "capital.loss", "hours.per.
↪week"]
categorical_features = [
    "workclass",
    "marital.status",
    "occupation",
    "relationship",
    "native.country",
]
ordinal_features = ["education"]
binary_features = ["sex"]
drop_features = ["race", "education.num"]
target_column = "income"
```

```
[51]: education_levels = [
    "Preschool",
    "1st-4th",
    "5th-6th",
    "7th-8th",
    "9th",
    "10th",
    "11th",
    "12th",
    "HS-grad",
    "Prof-school",
    "Assoc-voc",
    "Assoc-acdm",
    "Some-college",
    "Bachelors",
    "Masters",
    "Doctorate",
]
```

```
[52]: assert set(education_levels) == set(train_df["education"].unique())
```

```
[53]: numeric_transformer = make_pipeline(SimpleImputer(strategy="median"),
↳ StandardScaler())
tree_numeric_transformer = make_pipeline(SimpleImputer(strategy="median"))

categorical_transformer = make_pipeline(
    SimpleImputer(strategy="constant", fill_value="missing"),
    OneHotEncoder(handle_unknown="ignore"),
)

ordinal_transformer = make_pipeline(
    SimpleImputer(strategy="constant", fill_value="missing"),
    OrdinalEncoder(categories=[education_levels], dtype=int),
)

binary_transformer = make_pipeline(
    SimpleImputer(strategy="constant", fill_value="missing"),
    OneHotEncoder(drop="if_binary", dtype=int),
)

preprocessor = make_column_transformer(
    ("drop", drop_features),
    (numeric_transformer, numeric_features),
    (ordinal_transformer, ordinal_features),
    (binary_transformer, binary_features),
    (categorical_transformer, categorical_features),
)
```

```
[54]: X_train = train_df_nan.drop(columns=[target_column])
y_train = train_df_nan[target_column]

X_test = test_df_nan.drop(columns=[target_column])
y_test = test_df_nan[target_column]
```

#### 1.6.4 Do we have class imbalance?

- There is class imbalance. But without any context, **both classes seem equally important**.
- Let's use accuracy as our metric.

```
[55]: train_df_nan["income"].value_counts(normalize=True)
```

```
[55]: <=50K    0.757985
>50K      0.242015
Name: income, dtype: float64
```

```
[56]: scoring_metric = "accuracy"
```

```
[57]: import warnings

warnings.simplefilter(action="ignore", category=FutureWarning)
warnings.simplefilter(action="ignore", category=UserWarning)
```

Let's store all the results in a dictionary called `results`.

```
[58]: results = {}
```

```
[59]: from lightgbm.sklearn import LGBMClassifier
from xgboost import XGBClassifier

pipe_lr = make_pipeline(
    preprocessor, LogisticRegression(max_iter=2000, random_state=123)
)
pipe_rf = make_pipeline(preprocessor, RandomForestClassifier(random_state=123))
pipe_xgb = make_pipeline(
    preprocessor, XGBClassifier(random_state=123, eval_metric="logloss",
    ↪verbosity=0)
)
pipe_lgbm = make_pipeline(preprocessor, LGBMClassifier(random_state=123))
classifiers = {
    "logistic regression": pipe_lr,
    "random forest": pipe_rf,
    "XGBoost": pipe_xgb,
    "LightGBM": pipe_lgbm,
}
```

```
[60]: dummy = DummyClassifier(strategy="most_frequent")
results["Dummy"] = mean_std_cross_val_scores(
    dummy, X_train, y_train, return_train_score=True, scoring=scoring_metric
)
```

```
[61]: for (name, model) in classifiers.items():
    results[name] = mean_std_cross_val_scores(
        model, X_train, y_train, return_train_score=True, scoring=scoring_metric
    )
```

```
[62]: pd.DataFrame(results).T
```

```
[62]:
```





	fit_time	score_time	test_score \
Dummy	0.018 (+/- 0.004)	0.011 (+/- 0.001)	0.758 (+/- 0.000)
logistic regression	1.449 (+/- 0.057)	0.032 (+/- 0.006)	0.850 (+/- 0.006)
random forest	12.519 (+/- 0.380)	0.149 (+/- 0.012)	0.857 (+/- 0.004)
XGBoost	3.407 (+/- 2.349)	0.091 (+/- 0.027)	0.871 (+/- 0.004)
LightGBM	0.405 (+/- 0.061)	0.078 (+/- 0.003)	0.871 (+/- 0.004)

```
train_score
```



Dummy	0.758 (+/- 0.000)
logistic regression	0.851 (+/- 0.001)
random forest	1.000 (+/- 0.000)
XGBoost	0.908 (+/- 0.001)
LightGBM	0.892 (+/- 0.000)

- One problem is that often simple models are interpretable but not accurate.
- But more complex models (e.g., LightGBM) are less interpretable.

	Interpretable	Accurate
Complex model		
Simple model		

[Source](#)

```
[63]: # Reactivate default warning settings.
# As a best practice, try not to suppress warnings without a good reason
# and be sure to reactivate them when suppression is not needed anymore.

warnings.simplefilter(action="default", category=FutureWarning)
warnings.simplefilter(action="default", category=UserWarning)
```

### 1.6.5 Feature importances in linear models

- **Simpler** models are often **more interpretable** but **less accurate**.

Let's create and fit a pipeline with preprocessor and **logistic regression**.

```
[64]: pipe_lr = make_pipeline(preprocessor, LogisticRegression(max_iter=2000,
↳ random_state=2))
pipe_lr.fit(X_train, y_train);
```

```
[65]: ohe_feature_names = (
    pipe_rf.named_steps["columntransformer"]
    .named_transformers_["pipeline-4"]
    .named_steps["onehotencoder"]
    .get_feature_names_out(categorical_features)
    .tolist()
)
```

```
feature_names = (
    numeric_features + ordinal_features + binary_features + ohe_feature_names
)
pd.DataFrame(columns=feature_names)
```

[65]: Empty DataFrame

```
Columns: [age, fnlwgt, capital.gain, capital.loss, hours.per.week, education,
sex, workclass_Federal-gov, workclass_Local-gov, workclass_Never-worked,
workclass_Private, workclass_Self-emp-inc, workclass_Self-emp-not-inc,
workclass_State-gov, workclass_Without-pay, workclass_missing,
marital.status_Divorced, marital.status_Married-AF-spouse,
marital.status_Married-civ-spouse, marital.status_Married-spouse-absent,
marital.status_Never-married, marital.status_Separated, marital.status_Widowed,
occupation_Adm-clerical, occupation_Armed-Forces, occupation_Craft-repair,
occupation_Exec-managerial, occupation_Farming-fishing, occupation_Handlers-
cleaners, occupation_Machine-op-inspct, occupation_Other-service,
occupation_Priv-house-serv, occupation_Prof-specialty, occupation_Protective-
serv, occupation_Sales, occupation_Tech-support, occupation_Transport-moving,
occupation_missing, relationship_Husband, relationship_Not-in-family,
relationship_Other-relative, relationship_Own-child, relationship_Unmarried,
relationship_Wife, native.country_Cambodia, native.country_Canada,
native.country_China, native.country_Columbia, native.country_Cuba,
native.country_Dominican-Republic, native.country_Ecuador, native.country_El-
Salvador, native.country_England, native.country_France, native.country_Germany,
native.country_Greece, native.country_Guatemala, native.country_Haiti,
native.country_Holand-Netherlands, native.country_Honduras, native.country_Hong,
native.country_Hungary, native.country_India, native.country_Iran,
native.country_Ireland, native.country_Italy, native.country_Jamaica,
native.country_Japan, native.country_Laos, native.country_Mexico,
native.country_Nicaragua, native.country_Outlying-US(Guam-USVI-etc),
native.country_Peru, native.country_Philippines, native.country_Poland,
native.country_Portugal, native.country_Puerto-Rico, native.country_Scotland,
native.country_South, native.country_Taiwan, native.country_Thailand,
native.country_Trinidad&Tobago, native.country_United-States,
native.country_Vietnam, native.country_Yugoslavia, native.country_missing]
Index: []
```

[0 rows x 86 columns]

```
[66]: data = {
    "coefficient": pipe_lr.named_steps["logisticregression"].coef_[0].tolist(),
    "magnitude": np.absolute(
        pipe_lr.named_steps["logisticregression"].coef_[0].tolist()
    ),
}
coef_df = pd.DataFrame(data, index=feature_names).sort_values(
    "magnitude", ascending=False
```

```
)
coef_df.head()
```

```
[66]:
```

	coefficient	magnitude
capital.gain	2.355927	2.355927
marital.status_Married-AF-spouse	1.754646	1.754646
occupation_Priv-house-serv	-1.436944	1.436944
marital.status_Married-civ-spouse	1.341062	1.341062
relationship_Wife	1.274917	1.274917

- Increasing `capital.gain` is
  - likely to push the prediction towards “>50k” income class.
- Whereas `occupation_Priv-house-serv` is
  - likely to push the prediction towards “<=50K” income.

Can we get feature importances for non-linear models?

## 1.7 Model interpretability beyond linear models

We will be looking at three ways for model interpretability.

- `sklearn feature_importances_`
- [eli5](#) (stands for “explain like I’m 5”)
- [SHAP](#)

### 1.7.1 `sklearn feature_importances_`

- Many `sklearn` models have `feature_importances_` attribute.
- For **tree-based models**, it’s calculated based on **impurity** (gini index or information gain).
- For example, let’s look at `feature_importances_` of `RandomForestClassifier`.

Let’s create and fit a pipeline with preprocessor and random forest.

```
[67]: pipe_rf = make_pipeline(preprocessor, RandomForestClassifier(random_state=2))
pipe_rf.fit(X_train, y_train);
```

Which features are driving the predictions the most?

```
[68]: data = {
    "Importance": pipe_rf.named_steps["randomforestclassifier"].
        feature_importances_,
}
imps = pd.DataFrame(data=data, index=feature_names,).sort_values(
    by="Importance", ascending=False
)
imps.head(10)
```

```
[68]:
```

	Importance
fnlwgt	0.169580

age	0.153339
education	0.102953
capital.gain	0.097686
hours.per.week	0.085583
marital.status_Married-civ-spouse	0.064646
relationship_Husband	0.048896
capital.loss	0.033387
marital.status_Never-married	0.028629
occupation_Exec-managerial	0.020458

### 1.7.2 Key point

- Unlike the linear model coefficients, `feature_importances_` **do not have a sign!**
  - They tell us about **importance**, but *not* an “up or down”.
  - Indeed, increasing a feature may cause the prediction to first go up, and then go down.
  - This cannot happen in linear models, because they are linear.

Do these importances match with importances identified by logistic regression?

```
[69]: data = {
        "random forest importance":
            pipe_rf.named_steps["randomforestclassifier"].feature_importances_,
        "logistic regression importance":
            pipe_lr.named_steps["logisticregression"].coef_[0],
    }
    imps = pd.DataFrame(data=data, index=feature_names)
```

```
[70]: imps.sort_values(by="random forest importance", ascending=False).head()
```

```
[70]:
```

	random forest importance	logistic regression importance
fnlwgt	0.169580	0.078255
age	0.153339	0.359699
education	0.102953	0.184117
capital.gain	0.097686	2.355927
hours.per.week	0.085583	0.370219

Let's compare their top ten important feature lists:

```
[71]: col_rf = "random forest importance"
        col_lr = "logistic regression importance"

        ranking = pd.DataFrame({
            col_rf: imps[col_rf].sort_values(ascending=False).index,
            col_lr: imps[col_lr].sort_values(ascending=False, key=np.abs).index
        }).rename_axis('ranking')

        ranking.head(10)
```

```
[71]:
```

	random forest importance	logistic regression importance
ranking		
0	fnlwgt	capital.gain
1	age	marital.status_Married-AF-spouse
2	education	occupation_Priv-house-serv
3	capital.gain	marital.status_Married-civ-spouse
4	hours.per.week	relationship_Wife
5	marital.status_Married-civ-spouse	native.country_Columbia
6	relationship_Husband	occupation_Prof-specialty
7	capital.loss	occupation_Exec-managerial
8	marital.status_Never-married	native.country_Dominican-Republic
9	occupation_Exec-managerial	relationship_Own-child

- In their top 10 lists, both models agree on:
  - capital.gain
  - marital.status\_Married-civ-spouse
  - occupation\_Exec-managerial
- The actual numbers for random forests and logistic regression are not really comparable.

### 1.7.3 How can we get feature importances for non sklearn models?

- One way to do it is by using a tool called [eli5](#).

You'll have to install it

```
conda install -n cpsc330 -c conda-forge eli5
```

Let's look at feature importances for XGBClassifier.

```
[72]: import eli5

pipe_xgb = make_pipeline(
    preprocessor,
    XGBClassifier(random_state=123, eval_metric="logloss", verbosity=0))

warnings.simplefilter(action="ignore", category=UserWarning) # ignore warnings
pipe_xgb.fit(X_train, y_train);
warnings.simplefilter(action="default", category=UserWarning) # reactivate
↳ warnings

eli5_xgb = eli5.explain_weights(pipe_xgb.named_steps["xgbclassifier"],
    ↳ feature_names=feature_names)
eli5_xgb
```

```
[72]: Explanation(estimator="XGBClassifier(base_score=0.5, booster='gbtree',
colsample_bylevel=1,\n                colsample_bynode=1, colsample_bytree=1,
enable_categorical=False,\n                eval_metric='logloss', gamma=0,
gpu_id=-1, importance_type=None,\n                interaction_constraints='',
learning_rate=0.300000012,\n                max_delta_step=0, max_depth=6,
```

```

min_child_weight=1, missing=nan,\n                    monotone_constraints='()',
n_estimators=100, n_jobs=8,\n                    num_parallel_tree=1,
predictor='auto', random_state=123,\n                    reg_alpha=0, reg_lambda=1,
scale_pos_weight=1, subsample=1,\n                    tree_method='exact',
validate_parameters=1, verbosity=0)", description='\nXGBoost feature
importances; values are numbers 0 <= x <= 1;\nall values sum to 1.\n',
error=None, method='feature importances', is_regression=False, targets=None, fea
ture_importances=FeatureImportances(importances=[FeatureWeight(feature='marital.
status_Married-civ-spouse', weight=0.40608603, std=None, value=None),
FeatureWeight(feature='capital.gain', weight=0.054722264, std=None, value=None),
FeatureWeight(feature='relationship_Own-child', weight=0.044089068, std=None,
value=None), FeatureWeight(feature='education', weight=0.034879886, std=None,
value=None), FeatureWeight(feature='occupation_Other-service',
weight=0.032538496, std=None, value=None), FeatureWeight(feature='capital.loss',
weight=0.026842514, std=None, value=None),
FeatureWeight(feature='occupation_Prof-specialty', weight=0.024732435, std=None,
value=None), FeatureWeight(feature='occupation_Exec-managerial',
weight=0.01791228, std=None, value=None),
FeatureWeight(feature='occupation_Tech-support', weight=0.017841721, std=None,
value=None), FeatureWeight(feature='occupation_Handlers-cleaners',
weight=0.017215686, std=None, value=None),
FeatureWeight(feature='occupation_Machine-op-inspct', weight=0.016416635,
std=None, value=None), FeatureWeight(feature='occupation_Farming-fishing',
weight=0.016381254, std=None, value=None),
FeatureWeight(feature='workclass_Federal-gov', weight=0.015793154, std=None,
value=None), FeatureWeight(feature='age', weight=0.011652168, std=None,
value=None), FeatureWeight(feature='workclass_Self-emp-inc', weight=0.01078908,
std=None, value=None), FeatureWeight(feature='hours.per.week',
weight=0.0107445745, std=None, value=None),
FeatureWeight(feature='relationship_Wife', weight=0.010192984, std=None,
value=None), FeatureWeight(feature='sex', weight=0.010090632, std=None,
value=None), FeatureWeight(feature='relationship_Not-in-family',
weight=0.009422912, std=None, value=None),
FeatureWeight(feature='workclass_Self-emp-not-inc', weight=0.009076657,
std=None, value=None)], remaining=66), decision_tree=None,
highlight_spaces=None, transition_features=None, image=None)

```

Let's look at feature importances for LGBMClassifier.

```

[73]: pipe_lgbm = make_pipeline(preprocessor, LGBMClassifier(random_state=123))
pipe_lgbm.fit(X_train, y_train)
eli5_lgbm = eli5.explain_weights(
    pipe_lgbm.named_steps["lgbmclassifier"], feature_names=feature_names
)
eli5_lgbm

```

```
[73]: Explanation(estimator='LGBMClassifier(random_state=123)',
description='\nLightGBM feature importances; values are numbers 0 <= x <=
1;\nall values sum to 1.\n', error=None, method='feature importances',
is_regression=False, targets=None, feature_importances=FeatureImportances(import
ances=[FeatureWeight(feature='marital.status_Married-civ-spouse',
weight=0.35584397468549844, std=None, value=None),
FeatureWeight(feature='capital.gain', weight=0.19098036725150688, std=None,
value=None), FeatureWeight(feature='education', weight=0.13630962171648306,
std=None, value=None), FeatureWeight(feature='age', weight=0.08515574639356348,
std=None, value=None), FeatureWeight(feature='capital.loss',
weight=0.06393026262311322, std=None, value=None),
FeatureWeight(feature='hours.per.week', weight=0.0418456004162135, std=None,
value=None), FeatureWeight(feature='fnlwtg', weight=0.02451337553136395,
std=None, value=None), FeatureWeight(feature='occupation_Exec-managerial',
weight=0.013429664556178146, std=None, value=None),
FeatureWeight(feature='occupation_Prof-specialty', weight=0.012015760716975882,
std=None, value=None), FeatureWeight(feature='occupation_Other-service',
weight=0.0066740803861996744, std=None, value=None),
FeatureWeight(feature='sex', weight=0.006525060558280043, std=None, value=None),
FeatureWeight(feature='relationship_Wife', weight=0.005453787951464236,
std=None, value=None), FeatureWeight(feature='workclass_Self-emp-not-inc',
weight=0.005364964149799726, std=None, value=None),
FeatureWeight(feature='occupation_Farming-fishing', weight=0.005168542153226924,
std=None, value=None), FeatureWeight(feature='relationship_Own-child',
weight=0.0046119894507038505, std=None, value=None),
FeatureWeight(feature='occupation_Tech-support', weight=0.0032505444073185675,
std=None, value=None), FeatureWeight(feature='occupation_Sales',
weight=0.0024604164249372976, std=None, value=None),
FeatureWeight(feature='workclass_Private', weight=0.002385110169021698,
std=None, value=None), FeatureWeight(feature='workclass_Federal-gov',
weight=0.002384806796571508, std=None, value=None),
FeatureWeight(feature='occupation_Handlers-cleaners',
weight=0.00228880851674175, std=None, value=None)], remaining=66),
decision_tree=None, highlight_spaces=None, transition_features=None, image=None)
```

You can also look at feature importances for `RandomForestClassifier`, which we have already trained above.

```
[74]: eli5_rf = eli5.explain_weights(
        pipe_rf.named_steps["randomforestclassifier"], feature_names=feature_names
    )
eli5_rf
```

```
[74]: Explanation(estimator='RandomForestClassifier(random_state=2)',
description='\nRandom forest feature importances; values are numbers 0 <= x <=
1;\nall values sum to 1.\n', error=None, method='feature importances',
is_regression=False, targets=None, feature_importances=FeatureImportances(import
ances=[FeatureWeight(feature='fnlwtg', weight=0.16958005428552844,
```

```

std=0.00562712608917515, value=None), FeatureWeight(feature='age',
weight=0.1533390215043909, std=0.019798601881412622, value=None),
FeatureWeight(feature='education', weight=0.10295283451436564,
std=0.017386002656557657, value=None), FeatureWeight(feature='capital.gain',
weight=0.09768586081972082, std=0.023939780703838818, value=None),
FeatureWeight(feature='hours.per.week', weight=0.08558272511107902,
std=0.012506242582114328, value=None),
FeatureWeight(feature='marital.status_Married-civ-spouse',
weight=0.06464573433266022, std=0.06927178573889378, value=None),
FeatureWeight(feature='relationship_Husband', weight=0.04889639212082628,
std=0.055850610267523165, value=None), FeatureWeight(feature='capital.loss',
weight=0.03338747186463071, std=0.007844378859870025, value=None),
FeatureWeight(feature='marital.status_Never-married',
weight=0.02862861716671859, std=0.03698238581587818, value=None),
FeatureWeight(feature='occupation_Exec-managerial', weight=0.0204579927038537,
std=0.01053650089346291, value=None), FeatureWeight(feature='occupation_Prof-
specialty', weight=0.019333429958487774, std=0.009345251361284262, value=None),
FeatureWeight(feature='sex', weight=0.011773608675662732,
std=0.011047957330415109, value=None),
FeatureWeight(feature='relationship_Wife', weight=0.010990609952383412,
std=0.01122621572509277, value=None), FeatureWeight(feature='workclass_Private',
weight=0.009378492963718118, std=0.0018793227079254462, value=None),
FeatureWeight(feature='relationship_Not-in-family', weight=0.009322842064467672,
std=0.012121966770941614, value=None), FeatureWeight(feature='workclass_Self-
emp-not-inc', weight=0.00797575339268079, std=0.0017752094852695698,
value=None), FeatureWeight(feature='occupation_Other-service',
weight=0.007805796269149016, std=0.0052126895400305875, value=None),
FeatureWeight(feature='workclass_Self-emp-inc', weight=0.006597991594898411,
std=0.0031876281997523723, value=None), FeatureWeight(feature='relationship_Own-
child', weight=0.006592011991669542, std=0.011971256496172504, value=None),
FeatureWeight(feature='native.country_United-States',
weight=0.00637699697609062, std=0.0012210456325934707, value=None)],
remaining=66), decision_tree=None, highlight_spaces=None,
transition_features=None, image=None)

```

Let's compare them with weights what we got with `sklearn feature_importances_`

```

[75]: data = {
        "Importance": pipe_rf.named_steps["randomforestclassifier"].
        ↳feature_importances_,
    }
    sk_feat_imp_rf = pd.DataFrame(data=data, index=feature_names,).sort_values(
        by="Importance", ascending=False
    )
    sk_feat_imp_rf.head(10)

```



```
[75]:
```

	Importance
fnlwgt	0.169580
age	0.153339
education	0.102953
capital.gain	0.097686
hours.per.week	0.085583
marital.status_Married-civ-spouse	0.064646
relationship_Husband	0.048896
capital.loss	0.033387
marital.status_Never-married	0.028629
occupation_Exec-managerial	0.020458

Let's see them all together

```
[76]: def eli5_features(explain_weights):
        return [f.feature for f in explain_weights.feature_importances.importances]

eli5_rows = len(eli5_features(eli5_xgb))

pd.DataFrame({
    "XGB eli5": eli5_features(eli5_xgb),
    "LGBM eli5": eli5_features(eli5_lgbm),
    "RandomForest eli5": eli5_features(eli5_xgb),
    "RandomForest sklearn": sk_feat_imp_rf.head(eli5_rows).index
}).rename_axis('ranking')
```

```
[76]:
```

	XGB eli5	LGBM eli5 \
ranking		
0	marital.status_Married-civ-spouse	marital.status_Married-civ-spouse
1	capital.gain	capital.gain
2	relationship_Own-child	education
3	education	age
4	occupation_Other-service	capital.loss
5	capital.loss	hours.per.week
6	occupation_Prof-specialty	fnlwgt
7	occupation_Exec-managerial	occupation_Exec-managerial
8	occupation_Tech-support	occupation_Prof-specialty
9	occupation_Handlers-cleaners	occupation_Other-service
10	occupation_Machine-op-inspct	sex
11	occupation_Farming-fishing	relationship_Wife
12	workclass_Federal-gov	workclass_Self-emp-not-inc
13	age	occupation_Farming-fishing
14	workclass_Self-emp-inc	relationship_Own-child
15	hours.per.week	occupation_Tech-support
16	relationship_Wife	occupation_Sales
17	sex	workclass_Private
18	relationship_Not-in-family	workclass_Federal-gov
19	workclass_Self-emp-not-inc	occupation_Handlers-cleaners

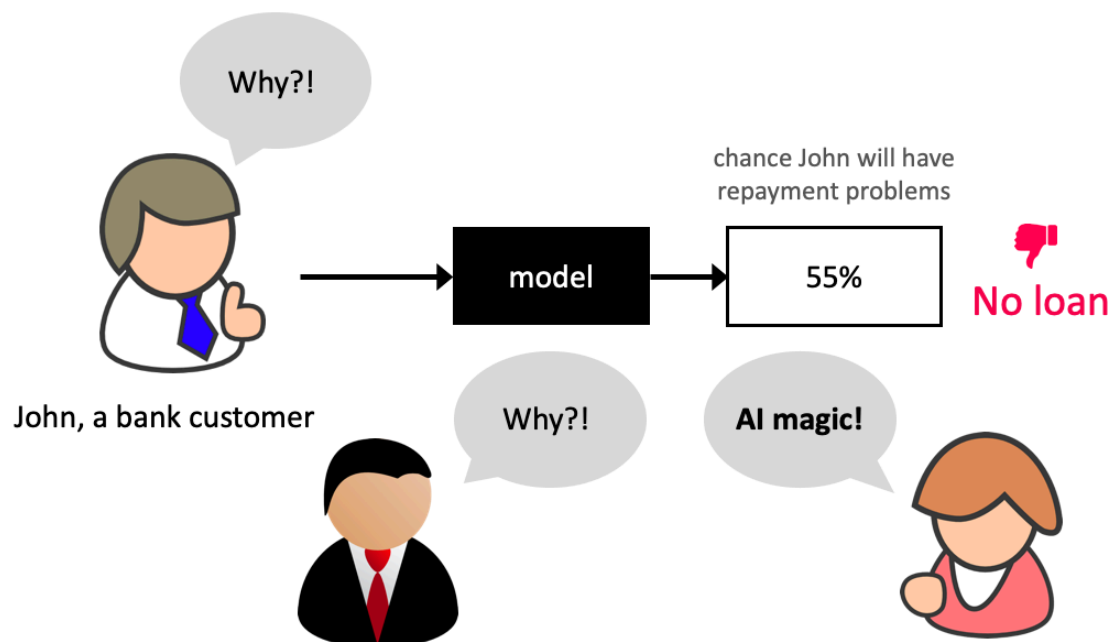
	RandomForest eli5	RandomForest sklearn
ranking		
0	marital.status_Married-civ-spouse	fnlwgt
1	capital.gain	age
2	relationship_Own-child	education
3	education	capital.gain
4	occupation_Other-service	hours.per.week
5	capital.loss	marital.status_Married-civ-spouse
6	occupation_Prof-specialty	relationship_Husband
7	occupation_Exec-managerial	capital.loss
8	occupation_Tech-support	marital.status_Never-married
9	occupation_Handlers-cleaners	occupation_Exec-managerial
10	occupation_Machine-op-inspct	occupation_Prof-specialty
11	occupation_Farming-fishing	sex
12	workclass_Federal-gov	relationship_Wife
13	age	workclass_Private
14	workclass_Self-emp-inc	relationship_Not-in-family
15	hours.per.week	workclass_Self-emp-not-inc
16	relationship_Wife	occupation_Other-service
17	sex	workclass_Self-emp-inc
18	relationship_Not-in-family	relationship_Own-child
19	workclass_Self-emp-not-inc	native.country_United-States

- These values tell us **globally** about which features are important
- But what if you want to explain a *specific* prediction?
- Some fancier tools can help us do this

## 1.8 SHAP (SHapley Additive exPlanations)

- A sophisticated *measure* of the **contribution of each feature**.
- [Lundberg and Lee, 2017](#)
- We won't go in details. You may refer to [Scott Lundberg's GitHub repo](#) if you are interested to know more.

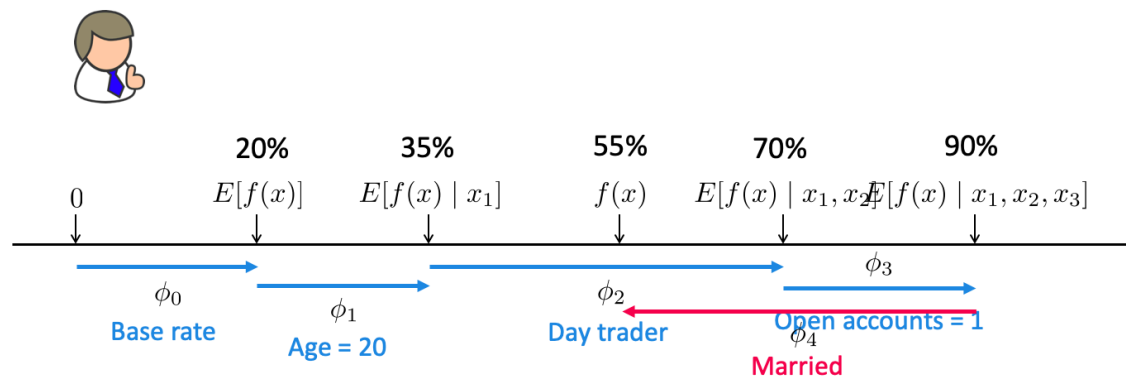
### 1.8.1 General idea



Source

### 1.8.2 General idea

- Provides following kind of explanation
  - Start at a **base rate** (e.g., how often people get their loans rejected).
  - Add one feature** at a time and see **how it impacts** the decision.



Source

Let's try it out on tree-based models.

First you'll have to install it.

```
conda install -n cpsc330 -c conda-forge shap
```

Let's create train and test dataframes with our transformed features.

```
[77]: X_train_enc = pd.DataFrame(
      data=preprocessor.transform(X_train).toarray(),
      columns=feature_names,
      index=X_train.index,
    )
X_train_enc.head()
```

```
[77]:
```

	age	fnlwgt	capital.gain	capital.loss	hours.per.week	\
5514	-0.921955	0.632531	-0.147166	-0.21768	-1.258387	
19777	-1.069150	-0.186155	-0.147166	-0.21768	-0.447517	
10781	-0.185975	-1.085437	-0.147166	-0.21768	-0.042081	
32240	-1.216346	-1.575119	-0.147166	-0.21768	-1.663822	
9876	-0.553965	1.588701	-0.147166	-0.21768	-0.042081	

	education	sex	workclass_Federal-gov	workclass_Local-gov	\
5514	8.0	1.0	0.0	0.0	
19777	8.0	0.0	0.0	0.0	
10781	13.0	0.0	0.0	0.0	
32240	12.0	0.0	0.0	0.0	
9876	13.0	1.0	0.0	1.0	

	workclass_Never-worked	...	native.country_Puerto-Rico	\
5514	0.0	...	0.0	
19777	0.0	...	0.0	
10781	0.0	...	0.0	
32240	0.0	...	0.0	
9876	0.0	...	0.0	

	native.country_Scotland	native.country_South	native.country_Taiwan	\
5514	0.0	0.0	0.0	
19777	0.0	0.0	0.0	
10781	0.0	0.0	0.0	
32240	0.0	0.0	0.0	
9876	0.0	0.0	0.0	

	native.country_Thailand	native.country_Trinidad&Tobago	\
5514	0.0	0.0	
19777	0.0	0.0	
10781	0.0	0.0	
32240	0.0	0.0	
9876	0.0	0.0	

	native.country_United-States	native.country_Vietnam	\
5514	1.0	0.0	
19777	1.0	0.0	
10781	1.0	0.0	
32240	1.0	0.0	

9876	1.0	0.0
	native.country_Yugoslavia	native.country_missing
5514	0.0	0.0
19777	0.0	0.0
10781	0.0	0.0
32240	0.0	0.0
9876	0.0	0.0

[5 rows x 86 columns]

```
[78]: X_test_enc = pd.DataFrame(
        data=preprocessor.transform(X_test).toarray(),
        columns=feature_names,
        index=X_test.index,
    )

X_test_enc.shape
```

[78]: (6513, 86)

Let's get SHAP values for train and test data.

```
[79]: import shap

lgbm_explainer = shap.TreeExplainer(pipe_lgbm.named_steps["lgbmclassifier"])
```

```
[80]: train_lgbm_shap_values = lgbm_explainer.shap_values(X_train_enc)
train_lgbm_shap_values[1].shape
```

LightGBM binary classifier with TreeExplainer shap values output has changed to a list of ndarray

[80]: (26048, 86)

```
[81]: test_lgbm_shap_values = lgbm_explainer.shap_values(X_test_enc)
test_lgbm_shap_values[1].shape
```

[81]: (6513, 86)

- For classification, it's a bit confusing. It gives SHAP arrays for both classes.
- Let's stick to shap values for class 1, i.e., income > 50K.

For each example and each feature we have a SHAP value.

```
[82]: train_lgbm_shap_values[1]
```

```
[82]: array([[ -4.23243013e-01,  -5.89878323e-02,  -2.65263112e-01, ...,
           9.63030623e-04,   0.00000000e+00,   5.74466631e-04],
          [ -6.83190014e-01,   1.15708200e-02,  -2.72482485e-01, ...,
```

```

8.17274476e-04, 0.00000000e+00, 8.09406158e-04],
[ 4.49106369e-01, -1.32455245e-01, -2.39454581e-01, ...,
 8.27603313e-04, 0.00000000e+00, 4.22023416e-03],
...,
[ 1.02714900e+00, 2.38119557e-02, -1.88163464e-01, ...,
 1.13580827e-03, 0.00000000e+00, 6.94390861e-04],
[ 6.37084418e-01, 2.90573592e-02, -3.03429292e-01, ...,
 9.70726909e-04, 0.00000000e+00, 2.16856964e-03],
[-1.24950883e+00, 1.19867799e-01, -2.23378846e-01, ...,
 9.70674774e-04, 0.00000000e+00, 9.73838044e-04]])

```

Let's look at the average SHAP values associated with each feature.

```

[83]: values = np.abs(train_lgbm_shap_values[1]).mean(0)
pd.DataFrame(data=values, index=feature_names, columns=["SHAP"]).sort_values(
    by="SHAP", ascending=False
).head(10)

```

```

[83]:
marital.status_Married-civ-spouse    SHAP    1.086269
age                                   0.823933
capital.gain                          0.572778
education                           0.409543
hours.per.week                      0.313901
sex                                  0.188874
capital.loss                        0.138607
relationship_Own-child              0.112871
occupation_Exec-managerial          0.107399
occupation_Prof-specialty           0.098181

```

- You can think of this as **global feature importance**
- But we'll see next that it gives you much more

## 1.9 SHAP plots

```

[84]: # load JS visualization code to notebook
shap.initjs()

```


<IPython.core.display.HTML object>

### 1.9.1 Dependence plot

```

[85]: shap.dependence_plot("age", train_lgbm_shap_values[1], X_train_enc)

```




The plot above shows effect of **age** feature on the prediction for class “>50K”.

- Each **dot** is a **single prediction** for examples above.
- The **x-axis** represents values of the **feature** age (scaled).
- The **y-axis** is the **SHAP value** for that feature
  - This represents how much knowing that **feature’s value changes the output** of the model for that example’s prediction.
- Lower values of age have smaller SHAP values for class “>50K”.
- Similarly, higher values of age also have a bit smaller SHAP values for class “>50K”, which makes sense.
- There is some optimal value of age between scaled age of 1 which gives highest SHAP values for for class “>50K”.
- Ignore the colour for now.
  - The colour corresponds to a **second feature (education** feature in this case) that may have an **interaction effect** with the feature we are plotting (age).

### 1.9.2 Summary plot

```
[86]: shap.summary_plot(train_lgbm_shap_values[1], X_train_enc)
```



The plot shows - (y-axis) The most important features for predicting the class - (x-axis) The direction of how feature values are going to drive the prediction. - (low feature values: blue; high feature values: red)

For example - Presence of the marital status of Married-civ-spouse seems to have bigger SHAP values for class 1 and absence seems to have smaller SHAP values for class 1. - Higher levels of education seem to have bigger SHAP values for class 1 whereas smaller levels of education have smaller SHAP values. - higher education pushes prediction towards >50K - lower education pushes prediction away from >50K

### 1.9.3 Force plot

- Let's try to explain predictions on a couple of examples from the test data.
- I'm sampling some examples where target is  $\leq 50K$  and some examples where target is  $> 50K$ .

```
[87]: 150k_indices, g50k_indices = y_test.reset_index().groupby('income').indices.  
      ↪ values()  
150k_indices, g50k_indices
```



```
[87]: (array([ 0, 1, 2, ..., 6508, 6509, 6511]),
      array([ 17, 18, 30, ..., 6505, 6510, 6512]))
```

```
[88]: ex_l50k_index = l50k_indices[10] # index of the tenth row with <=50K
      ex_g50k_index = g50k_indices[10] # index of the tenth row with >50K
      ex_l50k_index, ex_g50k_index
```

```
[88]: (10, 68)
```

See the rows with these indices:

```
[89]: y_test.iloc[[ex_l50k_index, ex_g50k_index]]
```

```
[89]: 345      <=50K
      23011     >50K
      Name: income, dtype: object
```

#### 1.9.4 Example with prediction <=50K

```
[90]: # pipe_lgbm.named_steps["lgbmclassifier"].
      ↪predict_proba(X_test_enc)[ex_l50k_index]
```


```
[91]: # pipe_lgbm.named_steps["lgbmclassifier"].predict(X_test_enc)[ex_l50k_index]
```

```
[92]: # pipe_lgbm.named_steps["lgbmclassifier"].predict(X_test_enc,
      ↪raw_score=True)[ex_l50k_index] # raw score of the model
```

```
[93]: # base_value
      lgbm_explainer.expected_value[1]
```

```
[93]: -2.3163172510079377
```

```
[94]: shap.force_plot(
      base_value=lgbm_explainer.expected_value[1],
      shap_values=test_lgbm_shap_values[1][ex_l50k_index, :],
      features=X_test_enc.iloc[ex_l50k_index, :],
      matplotlib=True,
      )
```



### 1.9.5 Example with prediction >50K

```
[95]: # pipe_lgbm.named_steps["lgbmclassifier"].  
      ↪ predict_proba(X_test_enc)[ex_g50k_index]
```


```
[96]: # pipe_lgbm.named_steps["lgbmclassifier"].predict(X_test_enc,   
      ↪ raw_score=True)[ex_g50k_index] # raw model score
```

```
[97]: # test_lgbm_shap_values[1][ex_g50k_index, :]
```

```
[98]: # base_value  
      lgbm_explainer.expected_value[1]
```

```
[98]: -2.3163172510079377
```

```
[99]: shap.force_plot(  
      base_value=lgbm_explainer.expected_value[1],  
      shap_values=test_lgbm_shap_values[1][ex_g50k_index, :],  
      features=X_test_enc.iloc[ex_g50k_index, :],  
      matplotlib=True,  
      )
```



Observations:

- Everything is with **respect to class 1** here.
- The base value for class 1 is -2.316. (You can think of this as the average raw score.)
- We see the forces that drive the prediction.
- That is, we can see the main factors pushing it from the base value (average over the dataset) to this particular prediction.
- Features that **push** the prediction to a **higher** value are shown in **red**.
- Features that **push** the prediction to a **lower** value are shown in **blue**.

**Note** A nice thing about SHAP values is that the feature importances sum to the prediction:

```
[100]: test_lgbm_shap_values[1][ex_g50k_index, :].sum() + lgbm_explainer.  
      ↪ expected_value[1]
```

```
[100]: 0.11096043410156309
```

```
[101]: # recall that
y_test.iloc[ex_g50k_index]
```

```
[101]: '>50K'
```

### 1.9.6 SHAP provides explainer for different kinds of models

- [TreeExplainer](#) (supports XGBoost, CatBoost, LightGBM)
- [DeepExplainer](#) (supports deep-learning models)
- [KernelExplainer](#) (supports kernel-based models)
- [GradientExplainer](#) (supports Keras and Tensorflow models)
- SHAP can also be used to explain text classification and image classification
- Example: In the picture below, red pixels represent positive SHAP values that increase the probability of the class, while blue pixels represent negative SHAP values that reduce the probability of the class.



[Source](#)

### 1.9.7 Other tools

- [lime](#) is another package.

### 1.9.8 In summary:

- So far we've only used sklearn models.
- Most sklearn models have some built-in measure of feature importances.
- On many tasks we need to move beyond sklearn, e.g. LightGBM, deep learning.
- These tools work on other models as well, which makes them extremely useful.

### 1.9.9 Why do we want this information?

Possible reasons:

- Identify features that are not useful and maybe remove them.
- Get guidance on what new data to collect.
  - New features related to useful features -> better results.
  - Don't bother collecting useless features -> save resources.
- Help explain why the model is making certain predictions.
  - Debugging, if the model is behaving strangely.
  - Regulatory requirements.
  - Fairness / bias.
  - Keep in mind this can be used on **deployment** predictions!

### 1.9.10 Questions for you

#### 1.9.11 True/False

1. You train a random forest on a binary classification problem with two classes [neg, pos]. A value of 0.580 for feat1 given by `feature_importances_` attribute of your model means that increasing the value of feat1 will drive us towards positive class. **FALSE** Feature importance ranges from 0 to 1. Never negative.
2. eli5 can be used to get feature importances for non `sklearn` models. **TRUE**
3. With SHAP you can only explain predictions on the training examples. **FALSE** It can also be used in deployment predictions.

[ ]: