

CPSC 330

Applied Machine Learning

1 Lecture 18: Time series

UBC 2022 Summer

Instructor: Mehrdad Oveis

1.1 Imports

```
[1]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from sklearn.compose import ColumnTransformer, make_column_transformer
from sklearn.dummy import DummyClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import (
    TimeSeriesSplit,
    cross_val_score,
    cross_validate,
    train_test_split,
)
from sklearn.pipeline import Pipeline, make_pipeline
from sklearn.preprocessing import OneHotEncoder, OrdinalEncoder, StandardScaler

plt.rcParams["font.size"] = 16
from datetime import datetime
```

1.2 Learning objectives

- Explain the pitfalls of train/test splitting with time series data.
- Appropriately split time series data, both train/test split and cross-validation.

- Perform time series feature engineering:
 - Encode time as various features in a tabular dataset
 - Create lag-based features
- Explain how can you forecast multiple time steps into the future.
- Explain the challenges of time series data with unequally spaced time points.
- At a high level, explain the concepts of seasonality and trends.

1.3 Motivation

- **Time series** is a collection of data points indexed in time order.
- Time series is everywhere:
 - Physical sciences (e.g., weather forecasting)
 - Economics, finance (e.g., stocks, market trends)
 - Engineering (e.g., energy consumption)
 - Social sciences
 - Sports analytics

Let's start with a simple example from [Introduction to Machine Learning with Python book](#).

In New York city there is a network of bike rental stations with a subscription system. The stations are all around the city. The anonymized data is available [here](#).

The task we will focus on is predicting how many people will rent a bicycle from a particular station for a given time and day. We might be interested in knowing this so that we know whether there will be any bikes left at the station for a particular day and time.

```
[2]: import mglearn

citibike = mglearn.datasets.load_citibike()
citibike.head()
```

```
[2]: starttime
2015-08-01 00:00:00      3
2015-08-01 03:00:00      0
2015-08-01 06:00:00      9
2015-08-01 09:00:00     41
2015-08-01 12:00:00     39
Freq: 3H, Name: one, dtype: int64
```

- The only feature we have is the date time feature.
 - Example: 2015-08-01 00:00:00
- The target is the number of rentals in the next 3 hours.
 - Example: 3 rentals between 2015-08-01 00:00:00 and 2015-08-01 03:00:00

```
[3]: citibike.index.min()
```


```
[3]: Timestamp('2015-08-01 00:00:00', freq='3H')
```

```
[4]: citibike.index.max()
```

```
[4]: Timestamp('2015-08-31 21:00:00', freq='3H')
```

We have data for August 2015.

```
[5]: plt.figure(figsize=(10, 3))
xticks = pd.date_range(start=citibike.index.min(), end=citibike.index.max(),
    ↪freq="D")
plt.xticks(xticks, xticks.strftime("%a %m-%d"), rotation=90, ha="left")
plt.plot(citibike, linewidth=2)
plt.xlabel("Date")
plt.ylabel("Rentals");
```



- We see the day and night pattern
- We see the weekend and weekday pattern
- Questions you might want to answer: How many people are likely to rent a bike at this station tomorrow at 3pm given everything we know about rentals in the past?
- We want to learn from the past and predict the future.

1.3.1 Train/test split for temporal data

- What will happen if we split this data the usual way?

```
[6]: train_df, test_df = train_test_split(citibike, test_size=0.2, random_state=123)
```

```
[7]: test_df.head()
```

```
[7]: starttime
2015-08-26 12:00:00    30
2015-08-12 09:00:00    10
2015-08-19 03:00:00     2
```

```
2015-08-07 12:00:00    22
2015-08-03 09:00:00     9
Name: one, dtype: int64
```

```
[8]: train_df.index.max()
```

```
[8]: Timestamp('2015-08-31 21:00:00')
```


```
[9]: test_df.index.min()
```

```
[9]: Timestamp('2015-08-01 12:00:00')
```

- So, we are training on data that came after our test data!
- If we want to forecast, **we aren't allowed to know what happened in the future!**
- There may be cases where this is OK, e.g. if you aren't trying to forecast and just want to understand your data (maybe you're not even splitting).
- But, for our purposes, we want to avoid this.

```
[10]: plt.figure(figsize=(10, 3))
train_df_sort = train_df.sort_index()
test_df_sort = test_df.sort_index()

plt.plot(train_df_sort, "b", label="train")
plt.plot(test_df_sort, "r", label="test")
plt.xticks(rotation="vertical")
plt.legend();
```



We'll split the data as follows:

- We have total 248 data points.

- We'll use the first 184 data points corresponding to the first 23 days as training data - And the remaining 64 data points corresponding to the remaining 8 days as test data.


```
[11]: citibike.shape
```

```
[11]: (248,)
```

```
[12]: n_train = 184
train_df = citibike[:184]
test_df = citibike[184:]
```

```
[13]: plt.figure(figsize=(10, 3))
train_df_sort = train_df.sort_index()
test_df_sort = test_df.sort_index()

plt.plot(train_df_sort, "b", label="train")
plt.plot(test_df_sort, "r", label="test")
plt.xticks(rotation="vertical")
plt.legend();
```



- This split is looking reasonable now.

1.3.2 Training models

- In this toy data, we just have a single feature: the date time feature.
- We need to encode this feature if we want to build machine learning models.
- A common way that dates are stored on computers is using POSIX time, which is the number of seconds since January 1970 00:00:00 (this is beginning of Unix time).
- Let's start with encoding this feature as a single integer representing this POSIX time.

```
[14]: X = (  
        citibike.index.astype("int64").values.reshape(-1, 1) // 10 ** 9  
    ) # convert to POSIX time by dividing by 10**9  
    y = citibike.values
```

```
[15]: y_train = train_df.values  
    y_test = test_df.values  
    # convert to POSIX time by dividing by 10**9  
    X_train = train_df.index.astype("int64").values.reshape(-1, 1) // 10 ** 9  
    X_test = test_df.index.astype("int64").values.reshape(-1, 1) // 10 ** 9
```

```
[16]: X_train[:10]
```

```
[16]: array([[1438387200],  
            [1438398000],  
            [1438408800],  
            [1438419600],  
            [1438430400],  
            [1438441200],  
            [1438452000],  
            [1438462800],  
            [1438473600],  
            [1438484400]])
```

```
[17]: y_train[:10]
```

```
[17]: array([ 3,  0,  9, 41, 39, 27, 12,  4,  3,  4])
```

- Our prediction task is a regression task.

Let's try random forest regression.

```
[18]: from sklearn.ensemble import RandomForestRegressor  
  
    regressor = RandomForestRegressor(n_estimators=100, random_state=0)  
    regressor.fit(X_train, y_train)  
  
    print("Train-set R^2: {:.2f}".format(regressor.score(X_train, y_train)))  
    print("Test-set R^2: {:.2f}".format(regressor.score(X_test, y_test)))
```

Train-set R²: 0.85

Test-set R²: -0.04

```
[19]: ## Code credit: https://learning.oreilly.com/library/view/introduction-to-machine/9781449369880/  
  
    def eval_on_features(features, target, regressor):  
        # split the given features into a training and a test set
```

```

X_train, X_test = features[:n_train], features[n_train:]
# also split the target array
y_train, y_test = target[:n_train], target[n_train:]
regressor.fit(X_train, y_train)
print("Train-set R^2: {:.2f}".format(regressor.score(X_train, y_train)))
print("Test-set R^2: {:.2f}".format(regressor.score(X_test, y_test)))
y_pred = regressor.predict(X_test)
y_pred_train = regressor.predict(X_train)
plt.figure(figsize=(10, 3))

plt.xticks(range(0, len(X), 8), xticks.strftime("%a %m-%d"), rotation=90,
↪ha="left")

plt.plot(range(n_train), y_train, label="train")
plt.plot(range(n_train, len(y_test) + n_train), y_test, "-", label="test")
plt.plot(range(n_train), y_pred_train, "--", label="prediction train")

plt.plot(
    range(n_train, len(y_test) + n_train), y_pred, "--", label="prediction_
↪test"
)
plt.legend(loc=(1.01, 0))
plt.xlabel("Date")
plt.ylabel("Rentals")

```

```
[20]: from sklearn.ensemble import RandomForestRegressor
```

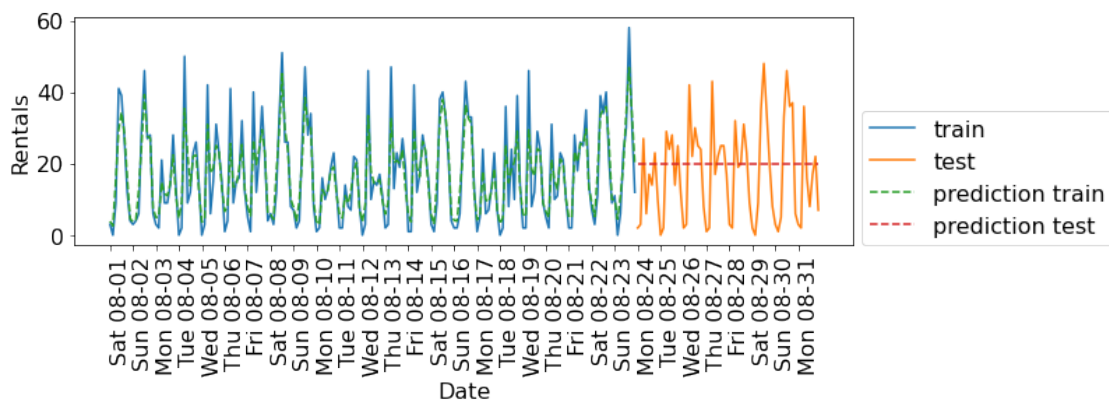
```

regressor = RandomForestRegressor(n_estimators=100, random_state=0)
eval_on_features(X, y, regressor)

```

Train-set R²: 0.85

Test-set R²: -0.04



- The predictions on the training score is pretty good
- But for the test data, a constant line is predicted ...
- What's going on?
- The model is based on only one feature: POSIX time feature.
- And the value of the POSIX time feature is outside the range of the feature values in the training set.
- Tree-based models cannot *extrapolate* to feature ranges outside the training data.
- The model predicted the target value of the closest point in the training set.

Can we come up with better features?

1.3.3 Feature engineering for date/time columns

- Note that our index is of this special type: `DatetimeIndex`. We can extract all kinds of interesting information from it.

```
[21]: citibike.index
```

```
[21]: DatetimeIndex(['2015-08-01 00:00:00', '2015-08-01 03:00:00',
                    '2015-08-01 06:00:00', '2015-08-01 09:00:00',
                    '2015-08-01 12:00:00', '2015-08-01 15:00:00',
                    '2015-08-01 18:00:00', '2015-08-01 21:00:00',
                    '2015-08-02 00:00:00', '2015-08-02 03:00:00',
                    ...
                    '2015-08-30 18:00:00', '2015-08-30 21:00:00',
                    '2015-08-31 00:00:00', '2015-08-31 03:00:00',
                    '2015-08-31 06:00:00', '2015-08-31 09:00:00',
                    '2015-08-31 12:00:00', '2015-08-31 15:00:00',
                    '2015-08-31 18:00:00', '2015-08-31 21:00:00'],
                    dtype='datetime64[ns]', name='starttime', length=248, freq='3H')
```

```
[22]: citibike.index.month_name()
```

```
[22]: Index(['August', 'August', 'August', 'August', 'August', 'August', 'August',
            'August', 'August', 'August',
            ...
            'August', 'August', 'August', 'August', 'August', 'August', 'August',
            'August', 'August', 'August'],
            dtype='object', name='starttime', length=248)
```

```
[23]: citibike.index.dayofweek
```

```
[23]: Int64Index([5, 5, 5, 5, 5, 5, 5, 5, 6, 6,
                ...
                6, 6, 0, 0, 0, 0, 0, 0, 0, 0],
                dtype='int64', name='starttime', length=248)
```

```
[24]: citibike.index.day_name()
```



```
[24]: Index(['Saturday', 'Saturday', 'Saturday', 'Saturday', 'Saturday', 'Saturday',
...
'Saturday', 'Saturday', 'Sunday', 'Sunday',
...
'Sunday', 'Sunday', 'Monday', 'Monday', 'Monday', 'Monday', 'Monday',
'Monday', 'Monday', 'Monday'],
dtype='object', name='starttime', length=248)
```

```
[25]: citibike.index.hour
```

```
[25]: Int64Index([ 0,  3,  6,  9, 12, 15, 18, 21,  0,  3,
...
18, 21,  0,  3,  6,  9, 12, 15, 18, 21],
dtype='int64', name='starttime', length=248)
```

- We noted before that the time of the day and day of the week seem quite important.
- Let's add these two features.

Let's first add the time of the day.


```
[26]: X_hour = citibike.index.hour.values.reshape(-1, 1)
X_hour[:10]
```

```
[26]: array([[ 0],
[ 3],
[ 6],
[ 9],
[12],
[15],
[18],
[21],
[ 0],
[ 3]])
```

```
[27]: eval_on_features(X_hour, y, regressor)
```

Train-set R^2 : 0.50

Test-set R^2 : 0.60



The scores are better than before.

Now let's add day of the week along with time of the day.

```
[28]: X_hour_week = np.hstack([
        citibike.index.dayofweek.values.reshape(-1, 1),
        citibike.index.hour.values.reshape(-1, 1),])


X_hour_week[:5]
```

```
[28]: array([[ 5,  0],
           [ 5,  3],
           [ 5,  6],
           [ 5,  9],
           [ 5, 12]])
```

```
[29]: eval_on_features(X_hour_week, y, regressor)
```

Train-set R^2 : 0.89

Test-set R^2 : 0.84



The results are much better. The time of the day and day of the week features are clearly helping.


- Do we need a complex model such as a random forest?
- Let's try Ridge with these features.

```
[30]: from sklearn.linear_model import Ridge

lr = Ridge()
eval_on_features(X_hour_week, y, lr)
```

Train-set R^2 : 0.16

Test-set R^2 : 0.13



- Ridge is performing **poorly** on the training as well as test data.
- It's not able to capture the periodic pattern.
- The reason is that we have encoded **time of day using integers**.
- A linear function can only learn a linear function of the time of day.
- What if we encode this feature as a **categorical** variable?

```
[31]: enc = OneHotEncoder()
      X_hour_week_onehot = enc.fit_transform(X_hour_week).toarray()
      X_hour_week_onehot.shape, X_hour_week.shape
```

```
[31]: ((248, 15), (248, 2))
```

```
[32]: enc.get_feature_names_out(['day', 'hour'])
```

```
[32]: array(['day_0', 'day_1', 'day_2', 'day_3', 'day_4', 'day_5', 'day_6',
            'hour_0', 'hour_3', 'hour_6', 'hour_9', 'hour_12', 'hour_15',
            'hour_18', 'hour_21'], dtype=object)
```

Or in a more readable format:

```
[33]: hour = ["%02d:00" % i for i in range(0, 24, 3)]
      day = ["Mon", "Tue", "Wed", "Thu", "Fri", "Sat", "Sun"]
      features_onehot = day + hour
      print(features_onehot)
```

```
['Mon', 'Tue', 'Wed', 'Thu', 'Fri', 'Sat', 'Sun', '00:00', '03:00', '06:00',
'09:00', '12:00', '15:00', '18:00', '21:00']
```

```
[34]: pd.DataFrame(X_hour_week_onehot, columns=features_onehot)
```

```
[34]:
```

	Mon	Tue	Wed	Thu	Fri	Sat	Sun	00:00	03:00	06:00	09:00	12:00	\
0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	1.0	0.0	0.0	

3	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	1.0	0.0
4	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	1.0
..
243	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0
244	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0
245	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
246	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
247	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0


	15:00	18:00	21:00
0	0.0	0.0	0.0
1	0.0	0.0	0.0
2	0.0	0.0	0.0
3	0.0	0.0	0.0
4	0.0	0.0	0.0
..
243	0.0	0.0	0.0
244	0.0	0.0	0.0
245	1.0	0.0	0.0
246	0.0	1.0	0.0
247	0.0	0.0	1.0

[248 rows x 15 columns]

```
[35]: eval_on_features(X_hour_week_onehot, y, Ridge())
```

Train-set R^2 : 0.53

Test-set R^2 : 0.62



- What if we add **interaction features**?
- We can do it using `sklearn`'s `PolynomialFeatures` transformer.

```
[36]: from sklearn.preprocessing import PolynomialFeatures
```

```
poly_transformer = PolynomialFeatures(
    degree=2, interaction_only=True, include_bias=False
)
```

```
[37]: toy = np.arange(12).reshape(-1, 3, order='F')
pd.DataFrame(toy, columns=['A', 'B', 'C',])
```

```
[37]:
```

	A	B	C
0	0	4	8
1	1	5	9
2	2	6	10
3	3	7	11

```
[38]: pd.DataFrame(poly_transformer.fit_transform(toy),
    columns=poly_transformer.get_feature_names_out(['A', 'B', 'C']))
```

```
[38]:
```

	A	B	C	A B	A C	B C
0	0.0	4.0	8.0	0.0	0.0	32.0
1	1.0	5.0	9.0	5.0	9.0	45.0
2	2.0	6.0	10.0	12.0	20.0	60.0
3	3.0	7.0	11.0	21.0	33.0	77.0

```
[39]: X_hour_week_onehot_poly = poly_transformer.fit_transform(X_hour_week_onehot)
pd.DataFrame(X_hour_week_onehot_poly).head()
```

```
[39]:
```


	0	1	2	3	4	5	6	7	8	9	...	110	111	112	113	\
0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0	0.0	...	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	1.0	0.0	...	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	1.0	...	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	

	114	115	116	117	118	119
0	0.0	0.0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0	0.0

[5 rows x 120 columns]

```
[40]: lr = Ridge()
eval_on_features(X_hour_week_onehot_poly, y, lr)
```

Train-set R^2 : 0.87
 Test-set R^2 : 0.85



```
[41]: X_hour_week_onehot_poly.shape, X_hour_week_onehot.shape, X_hour_week.shape
```

```
[41]: ((248, 120), (248, 15), (248, 2))
```

Let's see what are the column names for X_hour_week_onehot_poly

```
[42]: print(features_onehot)
```

```
['Mon', 'Tue', 'Wed', 'Thu', 'Fri', 'Sat', 'Sun', '00:00', '03:00', '06:00',
'09:00', '12:00', '15:00', '18:00', '21:00']
```

```
[43]: features_onehot_poly = poly_transformer.get_feature_names_out(features_onehot)
features_onehot_poly
```

```
[43]: array(['Mon', 'Tue', 'Wed', 'Thu', 'Fri', 'Sat', 'Sun', '00:00', '03:00',
'06:00', '09:00', '12:00', '15:00', '18:00', '21:00', 'Mon Tue',
'Mon Wed', 'Mon Thu', 'Mon Fri', 'Mon Sat', 'Mon Sun', 'Mon 00:00',
'Mon 03:00', 'Mon 06:00', 'Mon 09:00', 'Mon 12:00', 'Mon 15:00',
'Mon 18:00', 'Mon 21:00', 'Tue Wed', 'Tue Thu', 'Tue Fri',
'Tue Sat', 'Tue Sun', 'Tue 00:00', 'Tue 03:00', 'Tue 06:00',
'Tue 09:00', 'Tue 12:00', 'Tue 15:00', 'Tue 18:00', 'Tue 21:00',
'Wed Thu', 'Wed Fri', 'Wed Sat', 'Wed Sun', 'Wed 00:00',
'Wed 03:00', 'Wed 06:00', 'Wed 09:00', 'Wed 12:00', 'Wed 15:00',
'Wed 18:00', 'Wed 21:00', 'Thu Fri', 'Thu Sat', 'Thu Sun',
'Thu 00:00', 'Thu 03:00', 'Thu 06:00', 'Thu 09:00', 'Thu 12:00',
'Thu 15:00', 'Thu 18:00', 'Thu 21:00', 'Fri Sat', 'Fri Sun',
'Fri 00:00', 'Fri 03:00', 'Fri 06:00', 'Fri 09:00', 'Fri 12:00',
'Fri 15:00', 'Fri 18:00', 'Fri 21:00', 'Sat Sun', 'Sat 00:00',
'Sat 03:00', 'Sat 06:00', 'Sat 09:00', 'Sat 12:00', 'Sat 15:00',
'Sat 18:00', 'Sat 21:00', 'Sun 00:00', 'Sun 03:00', 'Sun 06:00',
'Sun 09:00', 'Sun 12:00', 'Sun 15:00', 'Sun 18:00', 'Sun 21:00',
'00:00 03:00', '00:00 06:00', '00:00 09:00', '00:00 12:00',
'00:00 15:00', '00:00 18:00', '00:00 21:00', '03:00 06:00',
'03:00 09:00', '03:00 12:00', '03:00 15:00', '03:00 18:00',
```

```

'03:00 21:00', '06:00 09:00', '06:00 12:00', '06:00 15:00',
'06:00 18:00', '06:00 21:00', '09:00 12:00', '09:00 15:00',
'09:00 18:00', '09:00 21:00', '12:00 15:00', '12:00 18:00',
'12:00 21:00', '15:00 18:00', '15:00 21:00', '18:00 21:00'],
dtype=object)

```

Let's examine the coefficients learned by Ridge. Note that many of the coefficients are zeros.

```
[44]: np.sum(lr.coef_ == 0), len(lr.coef_)
```

```
[44]: (49, 120)
```

```
[45]: np.array(features_onehot_poly)[lr.coef_ == 0]
```

```
[45]: array(['Mon Tue', 'Mon Wed', 'Mon Thu', 'Mon Fri', 'Mon Sat', 'Mon Sun',
'Tue Wed', 'Tue Thu', 'Tue Fri', 'Tue Sat', 'Tue Sun', 'Wed Thu',
'Wed Fri', 'Wed Sat', 'Wed Sun', 'Thu Fri', 'Thu Sat', 'Thu Sun',
'Fri Sat', 'Fri Sun', 'Sat Sun', '00:00 03:00', '00:00 06:00',
'00:00 09:00', '00:00 12:00', '00:00 15:00', '00:00 18:00',
'00:00 21:00', '03:00 06:00', '03:00 09:00', '03:00 12:00',
'03:00 15:00', '03:00 18:00', '03:00 21:00', '06:00 09:00',
'06:00 12:00', '06:00 15:00', '06:00 18:00', '06:00 21:00',
'09:00 12:00', '09:00 15:00', '09:00 18:00', '09:00 21:00',
'12:00 15:00', '12:00 18:00', '12:00 21:00', '15:00 18:00',
'15:00 21:00', '18:00 21:00'], dtype=object)
```

```
[46]: np.array(features_onehot_poly)[lr.coef_ != 0]
```

```
[46]: array(['Mon', 'Tue', 'Wed', 'Thu', 'Fri', 'Sat', 'Sun', '00:00', '03:00',
'06:00', '09:00', '12:00', '15:00', '18:00', '21:00', 'Mon 00:00',
'Mon 03:00', 'Mon 06:00', 'Mon 09:00', 'Mon 12:00', 'Mon 15:00',
'Mon 18:00', 'Mon 21:00', 'Tue 00:00', 'Tue 03:00', 'Tue 06:00',
'Tue 09:00', 'Tue 12:00', 'Tue 15:00', 'Tue 18:00', 'Tue 21:00',
'Wed 00:00', 'Wed 03:00', 'Wed 06:00', 'Wed 09:00', 'Wed 12:00',
'Wed 15:00', 'Wed 18:00', 'Wed 21:00', 'Thu 00:00', 'Thu 03:00',
'Thu 06:00', 'Thu 09:00', 'Thu 12:00', 'Thu 15:00', 'Thu 18:00',
'Thu 21:00', 'Fri 00:00', 'Fri 03:00', 'Fri 06:00', 'Fri 09:00',
'Fri 12:00', 'Fri 15:00', 'Fri 18:00', 'Fri 21:00', 'Sat 00:00',
'Sat 03:00', 'Sat 06:00', 'Sat 09:00', 'Sat 12:00', 'Sat 15:00',
'Sat 18:00', 'Sat 21:00', 'Sun 00:00', 'Sun 03:00', 'Sun 06:00',
'Sun 09:00', 'Sun 12:00', 'Sun 15:00', 'Sun 18:00', 'Sun 21:00'],
dtype=object)
```

```
[47]: features_nonzero = np.array(features_onehot_poly)[lr.coef_ != 0]
coef_nonzero = lr.coef_[lr.coef_ != 0]
```

```
[48]: pd.DataFrame(coef_nonzero, index=features_nonzero, columns=["Coefficient"]).
      ↪sort_values(
```

```
"Coefficient", ascending=False
)
```

```
[48]:
```

	Coefficient
Sat 09:00	15.196739
Wed 06:00	15.005809
Sat 12:00	13.437684
Sun 12:00	13.362009
Thu 06:00	10.907595
...	...
Sat 21:00	-6.085150
00:00	-11.693898
03:00	-12.111220
Sat 06:00	-13.757591
Sun 06:00	-18.033267

```
[71 rows x 1 columns]
```


- The coefficients make sense!
- If it's Saturday 09:00 or Wednesday 06:00, the model is likely to predict bigger number for rentals.
- If it's Midnight or 03:00 or Sunday 06:00, the model is likely to predict smaller number for rentals.

1.3.4 Cross-validation

What about cross-validation?

- We can't do regular cross-validation if we don't want to be predicting the past.
- If you carry out regular cross-validation, you'll be predicting the past given future which is not a realistic scenario for the deployment data.

```
[49]: mglearn.plots.plot_cross_validation()
```



There is `TimeSeriesSplit` for time series data.

```
[50]: from sklearn.model_selection import TimeSeriesSplit
```

```
[51]: X_toy = np.arange(0, 300, 10).reshape(-1, 2)
pd.DataFrame(X_toy)
```



```
[51]:      0    1
      0    0   10
      1   20   30
      2   40   50
      3   60   70
      4   80   90
      5  100  110
      6  120  130
      7  140  150
      8  160  170
      9  180  190
     10  200  210
     11  220  230
     12  240  250
     13  260  270
     14  280  290
```

```
[52]: tscv = TimeSeriesSplit(n_splits=4)
print("X_toy.shape:", X_toy.shape, "\n")
for train_idx, test_idx in tscv.split(X_toy):
    print("train_idx", train_idx)
    print("test_idx ", test_idx, "\n")
```

```
X_toy.shape: (15, 2)
```

```
train_idx [0 1 2]
test_idx  [3 4 5]
```

```
train_idx [0 1 2 3 4 5]
test_idx  [6 7 8]
```

```
train_idx [0 1 2 3 4 5 6 7 8]
test_idx  [ 9 10 11]
```

```
train_idx [ 0  1  2  3  4  5  6  7  8  9 10 11]
test_idx  [12 13 14]
```

Let's try it out with Ridge on the cities data.

```
[53]: lr = Ridge()
```

```
[54]: scores = cross_validate(
    lr, X_hour_week_onehot_poly, y, cv=TimeSeriesSplit(),
    ↪return_train_score=True
)
pd.DataFrame(scores)
```

```
[54]:
```

	fit_time	score_time	test_score	train_score
0	0.001307	0.000453	0.642676	0.873182
1	0.000588	0.000244	0.828405	0.874305
2	0.000705	0.000237	0.773851	0.901262
3	0.000749	0.000287	0.696712	0.889429
4	0.002546	0.000489	0.892733	0.863889

1.4 A more complicated dataset

[Rain in Australia](#) dataset. Predicting whether or not it will rain tomorrow based on today's measurements.

```
[55]: rain_df = pd.read_csv("data/weatherAUS.csv")
rain_df.head()
```

```
[55]:
```

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	\
0	2008-12-01	Albury	13.4	22.9	0.6	NaN	NaN	
1	2008-12-02	Albury	7.4	25.1	0.0	NaN	NaN	
2	2008-12-03	Albury	12.9	25.7	0.0	NaN	NaN	
3	2008-12-04	Albury	9.2	28.0	0.0	NaN	NaN	
4	2008-12-05	Albury	17.5	32.3	1.0	NaN	NaN	

	WindGustDir	WindGustSpeed	WindDir9am	...	Humidity9am	Humidity3pm	\
0	W	44.0	W	...	71.0	22.0	
1	WNW	44.0	NNW	...	44.0	25.0	
2	WSW	46.0	W	...	38.0	30.0	
3	NE	24.0	SE	...	45.0	16.0	
4	W	41.0	ENE	...	82.0	33.0	

	Pressure9am	Pressure3pm	Cloud9am	Cloud3pm	Temp9am	Temp3pm	RainToday	\
0	1007.7	1007.1	8.0	NaN	16.9	21.8	No	
1	1010.6	1007.8	NaN	NaN	17.2	24.3	No	
2	1007.6	1008.7	NaN	2.0	21.0	23.2	No	
3	1017.6	1012.8	NaN	NaN	18.1	26.5	No	
4	1010.8	1006.0	7.0	8.0	17.8	29.7	No	

	RainTomorrow
0	No
1	No
2	No
3	No
4	No

[5 rows x 23 columns]

```
[56]: rain_df.shape
```

```
[56]: (145460, 23)
```

Goals

- Can the date/time features help us predict the target value?
- Can we **forecast** into the future?

```
[57]: rain_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 145460 entries, 0 to 145459
Data columns (total 23 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Date                  145460 non-null object
 1   Location               145460 non-null object
 2   MinTemp                143975 non-null float64
 3   MaxTemp                144199 non-null float64
 4   Rainfall               142199 non-null float64
 5   Evaporation            82670 non-null  float64
 6   Sunshine               75625 non-null  float64
 7   WindGustDir            135134 non-null object
 8   WindGustSpeed          135197 non-null float64
 9   WindDir9am             134894 non-null object
10   WindDir3pm             141232 non-null object
11   WindSpeed9am           143693 non-null float64
12   WindSpeed3pm           142398 non-null float64
13   Humidity9am            142806 non-null float64
14   Humidity3pm            140953 non-null float64
15   Pressure9am            130395 non-null float64
16   Pressure3pm            130432 non-null float64
17   Cloud9am               89572 non-null  float64
18   Cloud3pm               86102 non-null  float64
19   Temp9am                143693 non-null float64
20   Temp3pm                141851 non-null float64
21   RainToday              142199 non-null object
22   RainTomorrow           142193 non-null object
dtypes: float64(16), object(7)
memory usage: 25.5+ MB
```

```
[58]: rain_df.describe(include="all")
```

```
[58]:
```

	Date	Location	MinTemp	MaxTemp	Rainfall	\
count	145460	145460	143975.000000	144199.000000	142199.000000	
unique	3436	49	NaN	NaN	NaN	
top	2013-11-12	Canberra	NaN	NaN	NaN	
freq	49	3436	NaN	NaN	NaN	
mean	NaN	NaN	12.194034	23.221348	2.360918	
std	NaN	NaN	6.398495	7.119049	8.478060	
min	NaN	NaN	-8.500000	-4.800000	0.000000	

25%	NaN	NaN	7.600000	17.900000	0.000000
50%	NaN	NaN	12.000000	22.600000	0.000000
75%	NaN	NaN	16.900000	28.200000	0.800000
max	NaN	NaN	33.900000	48.100000	371.000000

	Evaporation	Sunshine	WindGustDir	WindGustSpeed	WindDir9am	...	\
count	82670.000000	75625.000000	135134	135197.000000	134894	...	
unique	NaN	NaN	16	NaN	16	...	
top	NaN	NaN	W	NaN	N	...	
freq	NaN	NaN	9915	NaN	11758	...	
mean	5.468232	7.611178	NaN	40.035230	NaN	...	
std	4.193704	3.785483	NaN	13.607062	NaN	...	
min	0.000000	0.000000	NaN	6.000000	NaN	...	
25%	2.600000	4.800000	NaN	31.000000	NaN	...	
50%	4.800000	8.400000	NaN	39.000000	NaN	...	
75%	7.400000	10.600000	NaN	48.000000	NaN	...	
max	145.000000	14.500000	NaN	135.000000	NaN	...	

	Humidity9am	Humidity3pm	Pressure9am	Pressure3pm	\
count	142806.000000	140953.000000	130395.000000	130432.000000	
unique	NaN	NaN	NaN	NaN	
top	NaN	NaN	NaN	NaN	
freq	NaN	NaN	NaN	NaN	
mean	68.880831	51.539116	1017.64994	1015.255889	
std	19.029164	20.795902	7.10653	7.037414	
min	0.000000	0.000000	980.50000	977.100000	
25%	57.000000	37.000000	1012.90000	1010.400000	
50%	70.000000	52.000000	1017.60000	1015.200000	
75%	83.000000	66.000000	1022.40000	1020.000000	
max	100.000000	100.000000	1041.00000	1039.600000	

	Cloud9am	Cloud3pm	Temp9am	Temp3pm	RainToday	\
count	89572.000000	86102.000000	143693.000000	141851.00000	142199	
unique	NaN	NaN	NaN	NaN	2	
top	NaN	NaN	NaN	NaN	No	
freq	NaN	NaN	NaN	NaN	110319	
mean	4.447461	4.509930	16.990631	21.68339	NaN	
std	2.887159	2.720357	6.488753	6.93665	NaN	
min	0.000000	0.000000	-7.200000	-5.40000	NaN	
25%	1.000000	2.000000	12.300000	16.60000	NaN	
50%	5.000000	5.000000	16.700000	21.10000	NaN	
75%	7.000000	7.000000	21.600000	26.40000	NaN	
max	9.000000	9.000000	40.200000	46.70000	NaN	

	RainTomorrow
count	142193
unique	2

top	No
freq	110316
mean	NaN
std	NaN
min	NaN
25%	NaN
50%	NaN
75%	NaN
max	NaN

[11 rows x 23 columns]

- A number of missing values.
- Some target values are also missing. I'm dropping these rows.

```
[59]: rain_df = rain_df[rain_df["RainTomorrow"].notna()]
rain_df.shape
```

```
[59]: (142193, 23)
```

Parsing datetimes

- In general, datetimes are a huge pain! Think of all the formats: MM-DD-YY, DD-MM-YY, YY-MM-DD, MM/DD/YY, DD/MM/YY, DD/MM/YYYY, 20:45, 8:45am, 8:45 PM, 8:45a, 08:00, 8:10:20,
- No, seriously, dealing with datetimes is THE WORST.
 - Time zones.
 - Daylight savings...
- Thankfully, pandas does a pretty good job here.

```
[60]: dates_rain = pd.to_datetime(rain_df["Date"])
dates_rain
```

```
[60]: 0      2008-12-01
1      2008-12-02
2      2008-12-03
3      2008-12-04
4      2008-12-05
...
145454 2017-06-20
145455 2017-06-21
145456 2017-06-22
145457 2017-06-23
145458 2017-06-24
Name: Date, Length: 142193, dtype: datetime64[ns]
```

They are all the same format, so we can also compare dates:

```
[61]: dates_rain[1] - dates_rain[0]
```

```
[61]: Timedelta('1 days 00:00:00')
```

```
[62]: dates_rain[1] > dates_rain[0]
```

```
[62]: True
```

```
[63]: (dates_rain[1] - dates_rain[0]).total_seconds()
```

```
[63]: 86400.0
```

We can also easily extract information from the date columns.

```
[64]: dates_rain[1]
```

```
[64]: Timestamp('2008-12-02 00:00:00')
```

```
[65]: dates_rain[1].month_name()
```

```
[65]: 'December'
```

```
[66]: dates_rain[1].day_name()
```

```
[66]: 'Tuesday'
```

```
[67]: dates_rain[1].is_year_end
```

```
[67]: False
```

```
[68]: dates_rain[1].is_leap_year
```

```
[68]: True
```

```
[69]: dates_rain[dates_rain.map(lambda d: d.is_year_end and d.is_leap_year)].unique()
```

```
[69]: array(['2008-12-31T00:00:00.000000000', '2016-12-31T00:00:00.000000000'],  
      dtype='datetime64[ns]')
```

Above pandas identified the date column automatically. You can tell pandas to parse the dates when reading in the CSV:

```
[70]: rain_df = pd.read_csv("data/weatherAUS.csv", parse_dates=["Date"])  
      rain_df.head()
```

```
[70]:
```

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	\
0	2008-12-01	Albury	13.4	22.9	0.6	NaN	NaN	
1	2008-12-02	Albury	7.4	25.1	0.0	NaN	NaN	
2	2008-12-03	Albury	12.9	25.7	0.0	NaN	NaN	
3	2008-12-04	Albury	9.2	28.0	0.0	NaN	NaN	
4	2008-12-05	Albury	17.5	32.3	1.0	NaN	NaN	

	WindGustDir	WindGustSpeed	WindDir9am	...	Humidity9am	Humidity3pm	\
0	W	44.0	W	...	71.0	22.0	
1	WNW	44.0	NNW	...	44.0	25.0	
2	WSW	46.0	W	...	38.0	30.0	
3	NE	24.0	SE	...	45.0	16.0	
4	W	41.0	ENE	...	82.0	33.0	

	Pressure9am	Pressure3pm	Cloud9am	Cloud3pm	Temp9am	Temp3pm	RainToday	\
0	1007.7	1007.1	8.0	NaN	16.9	21.8	No	
1	1010.6	1007.8	NaN	NaN	17.2	24.3	No	
2	1007.6	1008.7	NaN	2.0	21.0	23.2	No	
3	1017.6	1012.8	NaN	NaN	18.1	26.5	No	
4	1010.8	1006.0	7.0	8.0	17.8	29.7	No	

	RainTomorrow
0	No
1	No
2	No
3	No
4	No

[5 rows x 23 columns]

```
[71]: rain_df["RainTomorrow"].isna().sum()
```

```
[71]: 3267
```

```
[72]: rain_df = rain_df[rain_df["RainTomorrow"].notna()]
rain_df.shape
```

```
[72]: (142193, 23)
```

```
[73]: rain_df["Date"].head()
```

```
[73]: 0    2008-12-01
1    2008-12-02
2    2008-12-03
3    2008-12-04
4    2008-12-05
Name: Date, dtype: datetime64[ns]
```

```
[74]: rain_df["Date"].unique().shape
```

```
[74]: (3436,)
```

1.5 Train/test splits

- Remember that we should not be calling the usual `train_test_split` with shuffling because

- If we want to forecast, we aren't allowed to know what happened in the future!

```
[75]: rain_df["Date"].min()
```

```
[75]: Timestamp('2007-11-01 00:00:00')
```

```
[76]: rain_df["Date"].max()
```

```
[76]: Timestamp('2017-06-25 00:00:00')
```

- It looks like we have 10 years of data.
- Let's use the last 2 years for test.

```
[77]: train_df = rain_df.query("Date <= 20150630")
test_df = rain_df.query("Date > 20150630")
```

```
[78]: len(train_df)
```

```
[78]: 107502
```

```
[79]: len(test_df)
```

```
[79]: 34691
```


```
[80]: len(test_df) / (len(train_df) + len(test_df))
```

```
[80]: 0.24397122221206388
```

As we can see, we're still using about 25% of our data as test data.

```
[81]: train_df_sort = train_df.query("Location == 'Sydney']").sort_values(by="Date")
test_df_sort = test_df.query("Location == 'Sydney']").sort_values(by="Date")

plt.plot(train_df_sort["Date"], train_df_sort["Rainfall"], "b", label="train")
plt.plot(test_df_sort["Date"], test_df_sort["Rainfall"], "r", label="test")
plt.xticks(rotation="vertical")
plt.legend()
plt.ylabel("Rainfall (mm)")
plt.title("Train/test rainfall in Sydney");
```

We're learning relationships from the blue part; predicting only using features in the red part from the day before.

Let's define a preprocessor with a column transformer.

```
[82]: train_df.columns
```

```
[82]: Index(['Date', 'Location', 'MinTemp', 'MaxTemp', 'Rainfall', 'Evaporation',
          'Sunshine', 'WindGustDir', 'WindGustSpeed', 'WindDir9am', 'WindDir3pm',
          'WindSpeed9am', 'WindSpeed3pm', 'Humidity9am', 'Humidity3pm',
          'Pressure9am', 'Pressure3pm', 'Cloud9am', 'Cloud3pm', 'Temp9am',
          'Temp3pm', 'RainToday', 'RainTomorrow'],
          dtype='object')
```

```
[83]: train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 107502 entries, 0 to 144733
Data columns (total 23 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Date            107502 non-null  datetime64[ns]
1   Location        107502 non-null  object
```

```

2  MinTemp      107050 non-null float64
3  MaxTemp      107292 non-null float64
4  Rainfall     106424 non-null float64
5  Evaporation  66221 non-null float64
6  Sunshine     62320 non-null float64
7  WindGustDir  100103 non-null object
8  WindGustSpeed 100146 non-null float64
9  WindDir9am   99515 non-null object
10 WindDir3pm   105314 non-null object
11 WindSpeed9am 106322 non-null float64
12 WindSpeed3pm 106319 non-null float64
13 Humidity9am  106112 non-null float64
14 Humidity3pm  106180 non-null float64
15 Pressure9am  97217 non-null float64
16 Pressure3pm  97253 non-null float64
17 Cloud9am     68523 non-null float64
18 Cloud3pm     67501 non-null float64
19 Temp9am      106705 non-null float64
20 Temp3pm      106816 non-null float64
21 RainToday    106424 non-null object
22 RainTomorrow 107502 non-null object
dtypes: datetime64[ns](1), float64(16), object(6)
memory usage: 19.7+ MB

```

- We have missing data.
- We have categorical features and numeric features.
- Let's define feature types.
- Let's start with **dropping the date** column and treating it as a **usual supervised machine learning** problem.

```

[84]: numeric_features = [
    "MinTemp",
    "MaxTemp",
    "Rainfall",
    "Evaporation",
    "Sunshine",
    "WindGustSpeed",
    "WindSpeed9am",
    "WindSpeed3pm",
    "Humidity9am",
    "Humidity3pm",
    "Pressure9am",
    "Pressure3pm",
    "Cloud9am",
    "Cloud3pm",
    "Temp9am",
    "Temp3pm",

```

```

]
categorical_features = [
    "Location",
    "WindGustDir",
    "WindDir9am",
    "WindDir3pm",
    "RainToday",
]
drop_features = [
    "Date",
    "RainTomorrow",
]

```

```

[85]: def preprocess_features(
    train_df,
    test_df,
    numeric_features,
    categorical_features,
    drop_features,
):

    all_features = set(numeric_features + categorical_features + drop_features)
    if set(train_df.columns) != all_features:
        print("Missing columns", set(train_df.columns) - all_features)
        print("Extra columns", all_features - set(train_df.columns))
        raise Exception("Columns do not match")

    numeric_transformer = make_pipeline(
        SimpleImputer(strategy="median"), StandardScaler()
    )
    categorical_transformer = make_pipeline(
        SimpleImputer(strategy="constant", fill_value="?"),
        OneHotEncoder(handle_unknown="ignore", sparse=False),
    )

    preprocessor = make_column_transformer(
        (numeric_transformer, numeric_features),
        (categorical_transformer, categorical_features),
        ("drop", drop_features),
    )
    preprocessor.fit(train_df)
    ohe_feature_names = (
        preprocessor.named_transformers_["pipeline-2"]
        .named_steps["onehotencoder"]
        .get_feature_names_out()
        .tolist()
    )

```

```

new_columns = numeric_features + ohe_feature_names

X_train_enc = pd.DataFrame(
    preprocessor.transform(train_df), index=train_df.index,
    columns=new_columns
)
X_test_enc = pd.DataFrame(
    preprocessor.transform(test_df), index=test_df.index,
    columns=new_columns
)

y_train = train_df["RainTomorrow"]
y_test = test_df["RainTomorrow"]

return X_train_enc, y_train, X_test_enc, y_test, preprocessor

```

```

[86]: X_train_enc, y_train, X_test_enc, y_test, preprocessor = preprocess_features(
    train_df,
    test_df,
    numeric_features,
    categorical_features,
    drop_features,
)

```

```

[87]: X_train_enc.head()

```

```

[87]:   MinTemp  MaxTemp  Rainfall  Evaporation  Sunshine  WindGustSpeed  \
0  0.204302 -0.027112 -0.205323   -0.140641  0.160729      0.298612
1 -0.741037  0.287031 -0.275008   -0.140641  0.160729      0.298612
2  0.125523  0.372706 -0.275008   -0.140641  0.160729      0.450132
3 -0.457435  0.701128 -0.275008   -0.140641  0.160729     -1.216596
4  0.850283  1.315134 -0.158867   -0.140641  0.160729      0.071330

   WindSpeed9am  WindSpeed3pm  Humidity9am  Humidity3pm  ...  x3_SE  x3_SSE  \
0      0.666166      0.599894      0.115428     -1.433514  ...    0.0    0.0
1     -1.125617      0.373275     -1.314929     -1.288002  ...    0.0    0.0
2      0.554180      0.826513     -1.632786     -1.045481  ...    0.0    0.0
3     -0.341712     -1.099749     -1.261953     -1.724539  ...    0.0    0.0
4     -0.789657      0.146656      0.698167     -0.899969  ...    0.0    0.0

   x3_SSW  x3_SW  x3_W  x3_WNW  x3_WSW  x4_?  x4_No  x4_Yes
0      0.0    0.0  0.0    1.0    0.0  0.0    1.0    0.0
1      0.0    0.0  0.0    0.0    1.0  0.0    1.0    0.0
2      0.0    0.0  0.0    0.0    1.0  0.0    1.0    0.0
3      0.0    0.0  0.0    0.0    0.0  0.0    1.0    0.0
4      0.0    0.0  0.0    0.0    0.0  0.0    1.0    0.0

```

[5 rows x 119 columns]

1.5.1 DummyClassifier

```
[88]: y_train.value_counts().to_frame().assign(ratio=y_train.  
      ↪value_counts(normalize=True))
```

```
[88]:      RainTomorrow      ratio  
No      83320  0.775055  
Yes     24182  0.224945
```

```
[89]: dc = DummyClassifier(strategy="prior")  
      dc.fit(train_df, y_train);
```

```
[90]: dc.score(train_df, y_train)
```

```
[90]: 0.7750553478074826
```

```
[91]: dc.score(test_df, y_test)
```

```
[91]: 0.7781845435415525
```

1.5.2 LogisticRegression

The function below trains a logistic regression model on the train set, reports train and test scores, and returns learned coefficients as a dataframe.

```
[92]: def score_lr_print_coef(preprocessor, train_df, y_train, test_df, y_test,   
      ↪columns):  
    lr_pipe = make_pipeline(preprocessor, LogisticRegression(max_iter=1000))  
    lr_pipe.fit(train_df, y_train)  
    print("Train score: {:.2f}".format(lr_pipe.score(train_df, y_train)))  
    print("Test score: {:.2f}".format(lr_pipe.score(test_df, y_test)))  
    lr_coef = pd.DataFrame(  
        data=lr_pipe.named_steps["logisticregression"].coef_.flatten(),  
        index=columns,  
        columns=["Coef"],  
    )  
    return lr_coef.sort_values(by="Coef", ascending=False)
```

```
[93]: score_lr_print_coef(preprocessor, train_df, y_train, test_df, y_test,   
      ↪X_train_enc.columns)
```

Train score: 0.85

Test score: 0.84

```
[93]:      Coef  
Humidity3pm  1.243181
```

```

x4_?          0.924002
Pressure9am    0.865490
x0_Witchcliffe 0.729016
WindGustSpeed  0.720411
...
x0_Townsville -0.718907
x0_Katherine  -0.725970
x0_Wollongong -0.749053
x0_MountGinini -0.965140
Pressure3pm    -1.221811

```

```
[119 rows x 1 columns]
```

1.5.3 Cross-validation

- We can carry out cross-validation using [TimeSeriesSplit](#).
- However, things are actually more complicated here because this dataset has **multiple time series**, one per location.

```
[94]: train_df
```

```

[94]:      Date Location  MinTemp  MaxTemp  Rainfall  Evaporation  Sunshine  \
0      2008-12-01  Albury    13.4    22.9        0.6           NaN        NaN
1      2008-12-02  Albury     7.4    25.1        0.0           NaN        NaN
2      2008-12-03  Albury    12.9    25.7        0.0           NaN        NaN
3      2008-12-04  Albury     9.2    28.0        0.0           NaN        NaN
4      2008-12-05  Albury    17.5    32.3        1.0           NaN        NaN
...      ...      ...      ...      ...      ...      ...      ...
144729 2015-06-26    Uluru     3.8    18.3        0.0           NaN        NaN
144730 2015-06-27    Uluru     2.5    17.1        0.0           NaN        NaN
144731 2015-06-28    Uluru     4.5    19.6        0.0           NaN        NaN
144732 2015-06-29    Uluru     7.6    22.0        0.0           NaN        NaN
144733 2015-06-30    Uluru     6.8    21.1        0.0           NaN        NaN

```

```

      WindGustDir  WindGustSpeed  WindDir9am  ...  Humidity9am  Humidity3pm  \
0              W             44.0          W  ...        71.0         22.0
1             WNW             44.0         NNW  ...        44.0         25.0
2             WSW             46.0          W  ...        38.0         30.0
3              NE             24.0          SE  ...        45.0         16.0
4              W             41.0         ENE  ...        82.0         33.0
...      ...      ...      ...      ...      ...      ...
144729          E             39.0         ESE  ...        73.0         37.0
144730          E             41.0         ESE  ...        69.0         40.0
144731         ENE             35.0         ESE  ...        69.0         39.0
144732         ESE             33.0          SE  ...        67.0         37.0
144733         ESE             35.0         ESE  ...        81.0         35.0

```

```

      Pressure9am  Pressure3pm  Cloud9am  Cloud3pm  Temp9am  Temp3pm  \

```

0	1007.7	1007.1	8.0	NaN	16.9	21.8
1	1010.6	1007.8	NaN	NaN	17.2	24.3
2	1007.6	1008.7	NaN	2.0	21.0	23.2
3	1017.6	1012.8	NaN	NaN	18.1	26.5
4	1010.8	1006.0	7.0	8.0	17.8	29.7
...
144729	1031.5	1027.6	NaN	NaN	8.8	17.2
144730	1029.9	1026.0	NaN	NaN	7.0	15.7
144731	1028.7	1025.0	NaN	3.0	8.9	18.0
144732	1027.2	1023.8	6.0	7.0	11.7	21.5
144733	1028.6	1025.2	3.0	NaN	10.6	20.2

	RainToday	RainTomorrow
0	No	No
1	No	No
2	No	No
3	No	No
4	No	No
...
144729	No	No
144730	No	No
144731	No	No
144732	No	No
144733	No	No

[107502 rows x 23 columns]

```
[95]: train_df.groupby(["Date", "Location"]).size().unique()
```

```
[95]: array([1])
```

```
[96]: train_df.groupby("Location").size().unique()
```

```
[96]: array([2369, 2305, 2290, 2314, 2226, 2304, 2444, 2310, 2696, 2277, 2229,
2219, 2467, 2289, 2463, 847, 2306, 2013, 2283, 2281, 2165, 2182,
2255, 845, 2220, 2282, 2280, 2048, 2252, 2274, 2228, 2235, 2611,
2313, 2279, 807, 2118, 1888, 2240, 2267])
```

```
[97]: train_df.groupby("Date").size().unique()
```

```
[97]: array([ 1, 2, 8, 7, 22, 23, 45, 44, 46, 43, 42, 39, 41, 48, 49, 47, 40])
```

```
[98]: train_df.sort_values(by=["Date"])
```

```
[98]:
```

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation \
45587	2007-11-01	Canberra	8.0	24.3	0.0	3.4
45588	2007-11-02	Canberra	14.0	26.9	3.6	4.4
45589	2007-11-03	Canberra	13.7	23.4	3.6	5.8

45590	2007-11-04	Canberra	13.3	15.5	39.8	7.2
45591	2007-11-05	Canberra	7.6	16.1	2.8	5.6
...
57415	2015-06-30	Ballarat	-0.3	10.5	0.0	NaN
119911	2015-06-30	PerthAirport	10.1	23.5	0.0	3.2
60455	2015-06-30	Bendigo	0.3	11.4	0.0	NaN
66473	2015-06-30	MelbourneAirport	3.2	13.2	0.0	0.8
144733	2015-06-30	Uluru	6.8	21.1	0.0	NaN

	Sunshine	WindGustDir	WindGustSpeed	WindDir9am	...	Humidity9am	\
45587	6.3	NW	30.0	SW	...	68.0	
45588	9.7	ENE	39.0	E	...	80.0	
45589	3.3	NW	85.0	N	...	82.0	
45590	9.1	NW	54.0	WNW	...	62.0	
45591	10.6	SSE	50.0	SSE	...	68.0	
...	
57415	NaN	S	26.0	NaN	...	99.0	
119911	5.8	NNE	31.0	NE	...	48.0	
60455	NaN	W	19.0	NaN	...	89.0	
66473	3.9	N	20.0	N	...	91.0	
144733	NaN	ESE	35.0	ESE	...	81.0	

	Humidity3pm	Pressure9am	Pressure3pm	Cloud9am	Cloud3pm	Temp9am	\
45587	29.0	1019.7	1015.0	7.0	7.0	14.4	
45588	36.0	1012.4	1008.4	5.0	3.0	17.5	
45589	69.0	1009.5	1007.2	8.0	7.0	15.4	
45590	56.0	1005.5	1007.0	2.0	7.0	13.5	
45591	49.0	1018.3	1018.5	7.0	7.0	11.1	
...	
57415	63.0	1029.5	1027.7	NaN	8.0	4.7	
119911	33.0	1023.6	1021.7	7.0	6.0	13.3	
60455	56.0	1029.3	1027.4	8.0	7.0	6.4	
66473	50.0	1029.6	1027.3	2.0	7.0	5.3	
144733	35.0	1028.6	1025.2	3.0	NaN	10.6	

	Temp3pm	RainToday	RainTomorrow
45587	23.6	No	Yes
45588	25.7	Yes	Yes
45589	20.2	Yes	Yes
45590	14.1	Yes	Yes
45591	15.4	Yes	No
...
57415	9.3	No	No
119911	22.2	No	No
60455	10.5	No	No
66473	11.9	No	No
144733	20.2	No	No

[107502 rows x 23 columns]

- It seems the dataframe is sorted by location, and then time.
- Our first approach will be to ignore the fact that we have multiple time series and just **OHE the location**
- We'll have **multiple measurements for a given timestamp**, and that's OK.
- But, TimeSeriesSplit expects the dataframe to be sorted by date so...

```
[99]: train_df_ordered = train_df.sort_values(by=["Date"])
      y_train_ordered = train_df_ordered["RainTomorrow"]
```

```
[100]: train_df_ordered
```

```
[100]:
```

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	\
45587	2007-11-01	Canberra	8.0	24.3	0.0	3.4	
45588	2007-11-02	Canberra	14.0	26.9	3.6	4.4	
45589	2007-11-03	Canberra	13.7	23.4	3.6	5.8	
45590	2007-11-04	Canberra	13.3	15.5	39.8	7.2	
45591	2007-11-05	Canberra	7.6	16.1	2.8	5.6	
...	
57415	2015-06-30	Ballarat	-0.3	10.5	0.0	NaN	
119911	2015-06-30	PerthAirport	10.1	23.5	0.0	3.2	
60455	2015-06-30	Bendigo	0.3	11.4	0.0	NaN	
66473	2015-06-30	MelbourneAirport	3.2	13.2	0.0	0.8	
144733	2015-06-30	Uluru	6.8	21.1	0.0	NaN	

	Sunshine	WindGustDir	WindGustSpeed	WindDir9am	...	Humidity9am	\
45587	6.3	NW	30.0	SW	...	68.0	
45588	9.7	ENE	39.0	E	...	80.0	
45589	3.3	NW	85.0	N	...	82.0	
45590	9.1	NW	54.0	WNW	...	62.0	
45591	10.6	SSE	50.0	SSE	...	68.0	
...	
57415	NaN	S	26.0	NaN	...	99.0	
119911	5.8	NNE	31.0	NE	...	48.0	
60455	NaN	W	19.0	NaN	...	89.0	
66473	3.9	N	20.0	N	...	91.0	
144733	NaN	ESE	35.0	ESE	...	81.0	

	Humidity3pm	Pressure9am	Pressure3pm	Cloud9am	Cloud3pm	Temp9am	\
45587	29.0	1019.7	1015.0	7.0	7.0	14.4	
45588	36.0	1012.4	1008.4	5.0	3.0	17.5	
45589	69.0	1009.5	1007.2	8.0	7.0	15.4	
45590	56.0	1005.5	1007.0	2.0	7.0	13.5	
45591	49.0	1018.3	1018.5	7.0	7.0	11.1	
...	
57415	63.0	1029.5	1027.7	NaN	8.0	4.7	

119911	33.0	1023.6	1021.7	7.0	6.0	13.3
60455	56.0	1029.3	1027.4	8.0	7.0	6.4
66473	50.0	1029.6	1027.3	2.0	7.0	5.3
144733	35.0	1028.6	1025.2	3.0	NaN	10.6

	Temp3pm	RainToday	RainTomorrow
45587	23.6	No	Yes
45588	25.7	Yes	Yes
45589	20.2	Yes	Yes
45590	14.1	Yes	Yes
45591	15.4	Yes	No
...
57415	9.3	No	No
119911	22.2	No	No
60455	10.5	No	No
66473	11.9	No	No
144733	20.2	No	No

[107502 rows x 23 columns]

```
[101]: lr_pipe = make_pipeline(preprocessor, LogisticRegression(max_iter=1000))
cross_val_score(lr_pipe, train_df_ordered, y_train_ordered,
cv=TimeSeriesSplit()).mean()
```

[101]: 0.8478874811631412

```
[102]: train_df_sydney = train_df.query("Location == 'Sydney'").sort_values(by="Date")
cross_val_score(
    lr_pipe, train_df_sydney, train_df_sydney["RainTomorrow"],
cv=TimeSeriesSplit()).mean()
```

[102]: 0.8317241379310344

```
[103]: def cross_val_score_loc(train_df_loc):
        # print(train_df_loc['Location'].unique())
        return cross_val_score(
            lr_pipe, train_df_loc, train_df_loc["RainTomorrow"],
cv=TimeSeriesSplit()).mean()

location_results = train_df.groupby("Location").apply(cross_val_score_loc)
location_results.head(3)
```

```
[103]: Location
Adelaide    0.862944
Albany      0.804688
Albury      0.875066
dtype: float64
```

```
[104]: location_results.describe()
```

```
[104]: count    49.000000
      mean     0.853753
      std     0.043287
      min     0.776316
      25%     0.813008
      50%     0.856511
      75%     0.884156
      max     0.943536
      dtype: float64
```

```
[105]: location_results.rename("CV Score").sort_values(ascending=False).to_frame().
      ↪reset_index()
```

```
[105]:
```

	Location	CV Score
0	Woomera	0.943536
1	AliceSprings	0.934026
2	Uluru	0.928358
3	Mildura	0.921579
4	Perth	0.912895
5	Cobar	0.908707
6	PearceRAAF	0.908504
7	Moree	0.906667
8	PerthAirport	0.903158
9	Katherine	0.892199
10	WaggaWagga	0.890000
11	Townsville	0.888312
12	Bendigo	0.884156
13	Tuggeranong	0.880739
14	SalmonGums	0.879032
15	Albury	0.875066
16	Nhil	0.874286
17	Nuriootpa	0.872632
18	Richmond	0.872237
19	Witchcliffe	0.865416
20	BadgerysCreek	0.864151
21	Adelaide	0.862944
22	Canberra	0.862361
23	Penrith	0.861333
24	Brisbane	0.856511
25	Darwin	0.854501
26	GoldCoast	0.845144
27	Launceston	0.842708
28	Ballarat	0.839063
29	MelbourneAirport	0.833158
30	Sydney	0.831724

31	MountGambier	0.827013
32	SydneyAirport	0.824737
33	Watsonia	0.820000
34	Wollongong	0.817507
35	Melbourne	0.813134
36	Dartmoor	0.813008
37	NorahHead	0.809730
38	Walpole	0.809065
39	Sale	0.808947
40	Williamstown	0.808917
41	Hobart	0.808780
42	Cairns	0.807792
43	CoffsHarbour	0.805930
44	Albany	0.804688
45	MountGinini	0.803857
46	Newcastle	0.795733
47	Portland	0.783641
48	NorfolkIsland	0.776316

Just as practice, for now, let's ignore location and just consider date.

1.6 Encoding date/time as feature(s)

- Can we use the `Date` to help us predict the target?
- Probably! E.g. different amounts of rain in different seasons.
- This is feature engineering!

1.6.1 Encoding time as an number

- Idea 1: create a column of “days since Nov 1, 2007” which is the first day in the dataset.

```
[106]: train_df["Date"].min()
```

```
[106]: Timestamp('2007-11-01 00:00:00')
```

```
[107]: train_df = rain_df.query("Date <= 20150630")
       test_df = rain_df.query("Date > 20150630")
```

```
[108]: first_day = train_df["Date"].min()

       train_df = train_df.assign(
           Days_since=train_df["Date"].apply(lambda x: (x - first_day).days)
       )
       test_df = test_df.assign(
           Days_since=test_df["Date"].apply(lambda x: (x - first_day).days)
       )
```

```
[109]: train_df.sort_values(by="Date").head()
```

```
[109]:
```

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	\
45587	2007-11-01	Canberra	8.0	24.3	0.0	3.4	6.3	
45588	2007-11-02	Canberra	14.0	26.9	3.6	4.4	9.7	
45589	2007-11-03	Canberra	13.7	23.4	3.6	5.8	3.3	
45590	2007-11-04	Canberra	13.3	15.5	39.8	7.2	9.1	
45591	2007-11-05	Canberra	7.6	16.1	2.8	5.6	10.6	

	WindGustDir	WindGustSpeed	WindDir9am	...	Humidity3pm	Pressure9am	\
45587	NW	30.0	SW	...	29.0	1019.7	
45588	ENE	39.0	E	...	36.0	1012.4	
45589	NW	85.0	N	...	69.0	1009.5	
45590	NW	54.0	WNW	...	56.0	1005.5	
45591	SSE	50.0	SSE	...	49.0	1018.3	

	Pressure3pm	Cloud9am	Cloud3pm	Temp9am	Temp3pm	RainToday	\
45587	1015.0	7.0	7.0	14.4	23.6	No	
45588	1008.4	5.0	3.0	17.5	25.7	Yes	
45589	1007.2	8.0	7.0	15.4	20.2	Yes	
45590	1007.0	2.0	7.0	13.5	14.1	Yes	
45591	1018.5	7.0	7.0	11.1	15.4	Yes	

	RainTomorrow	Days_since
45587	Yes	0
45588	Yes	1
45589	Yes	2
45590	Yes	3
45591	No	4

[5 rows x 24 columns]

```
[110]: X_train_enc, y_train, X_test_enc, y_test, preprocessor = preprocess_features(
    train_df,
    test_df,
    numeric_features + ["Days_since"],
    categorical_features,
    drop_features,
)
```

```
[111]: score_lr_print_coeff(preprocessor, train_df, y_train, test_df, y_test,
    ↪X_train_enc.columns)
```

Train score: 0.85
Test score: 0.84

```
[111]:
```

	Coef
Humidity3pm	1.243092
x4_?	0.935320
Pressure9am	0.864500

```

x0_Witchcliffe 0.731007
WindGustSpeed 0.720028
...
x0_Townsville -0.716770
x0_Katherine -0.738735
x0_Wollongong -0.746094
x0_MountGinini -0.963654
Pressure3pm -1.221826

```

```
[120 rows x 1 columns]
```

- Not much improvement in the scores
- Can you think of other ways to generate features from the Date column?

1.6.2 One-hot encoding of the month

- Idea 2: month
- The month seems relevant here. How should we encode the month?
- Encode it as a categorical variable?

```
[112]: train_df = rain_df.query("Date <= 20150630")
test_df = rain_df.query("Date > 20150630")
```

```
[113]: # use month_name() to get the actual string
train_df = train_df.assign(Month=train_df["Date"].apply(lambda x: x.
    ↪month_name()))
test_df = test_df.assign(Month=test_df["Date"].apply(lambda x: x.month_name()))
```

```
[114]: # To ensure correct month ordering
months_cat = pd.CategoricalDtype(ordered=True,
    categories=['January', 'February', 'March', 'April', 'May', 'June',
    ↪'July', 'August', 'September', 'October', 'November', 'December'])

train_df["Month"] = train_df["Month"].astype(months_cat)
test_df["Month"] = test_df["Month"].astype(months_cat)
```

```
[115]: train_df[["Date", "Month"]].sort_values(by="Month")
```

```
[115]:
```

	Date	Month
32616	2015-01-04	January
38285	2014-01-20	January
38284	2014-01-19	January
38283	2014-01-18	January
38282	2014-01-17	January
...
47724	2013-12-05	December
47723	2013-12-04	December

```
47722 2013-12-03  December
47720 2013-12-01  December
0      2008-12-01  December
```

```
[107502 rows x 2 columns]
```

```
[116]: X_train_enc, y_train, X_test_enc, y_test, preprocessor = preprocess_features(
        train_df, test_df, numeric_features, categorical_features + ["Month"],
        drop_features
    )
```

```
[117]: score_lr_print_coeff(preprocessor, train_df, y_train, test_df, y_test,
        X_train_enc.columns)
```

```
Train score: 0.85
```

```
Test score: 0.84
```

```
[117]:
```

	Coef
Humidity3pm	1.266873
x4_?	0.942135
Pressure9am	0.799146
x0_Witchcliffe	0.748923
WindGustSpeed	0.705617
...	...
x0_Darwin	-0.736057
x0_Wollongong	-0.748181
x0_Townsville	-0.903216
x0_Katherine	-0.930322
Pressure3pm	-1.182015

```
[131 rows x 1 columns]
```

1.6.3 One-hot encoding seasons

How about just summer/winter as a feature?

```
[118]: def get_season(month):
        WINTER_MONTHS = ["June", "July", "August"]
        AUTUMN_MONTHS = ["March", "April", "May"]
        SUMMER_MONTHS = ["December", "January", "February"]
        SPRING_MONTHS = ["September", "October", "November"]
        if month in WINTER_MONTHS:
            return "Winter"
        elif month in AUTUMN_MONTHS:
            return "Autumn"
        elif month in SUMMER_MONTHS:
            return "Summer"
        else:
```

```
return "Fall"
```

```
[119]: train_df = train_df.assign(Season=train_df["Month"].apply(get_season))
test_df = test_df.assign(Season=test_df["Month"].apply(get_season))
```

```
[120]: train_df
```

```
[120]:
```

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	\
0	2008-12-01	Albury	13.4	22.9	0.6	NaN	NaN	
1	2008-12-02	Albury	7.4	25.1	0.0	NaN	NaN	
2	2008-12-03	Albury	12.9	25.7	0.0	NaN	NaN	
3	2008-12-04	Albury	9.2	28.0	0.0	NaN	NaN	
4	2008-12-05	Albury	17.5	32.3	1.0	NaN	NaN	
...	
144729	2015-06-26	Uluru	3.8	18.3	0.0	NaN	NaN	
144730	2015-06-27	Uluru	2.5	17.1	0.0	NaN	NaN	
144731	2015-06-28	Uluru	4.5	19.6	0.0	NaN	NaN	
144732	2015-06-29	Uluru	7.6	22.0	0.0	NaN	NaN	
144733	2015-06-30	Uluru	6.8	21.1	0.0	NaN	NaN	

	WindGustDir	WindGustSpeed	WindDir9am	...	Pressure9am	Pressure3pm	\
0	W	44.0	W	...	1007.7	1007.1	
1	WNW	44.0	NNW	...	1010.6	1007.8	
2	WSW	46.0	W	...	1007.6	1008.7	
3	NE	24.0	SE	...	1017.6	1012.8	
4	W	41.0	ENE	...	1010.8	1006.0	
...	
144729	E	39.0	ESE	...	1031.5	1027.6	
144730	E	41.0	ESE	...	1029.9	1026.0	
144731	ENE	35.0	ESE	...	1028.7	1025.0	
144732	ESE	33.0	SE	...	1027.2	1023.8	
144733	ESE	35.0	ESE	...	1028.6	1025.2	

	Cloud9am	Cloud3pm	Temp9am	Temp3pm	RainToday	RainTomorrow	\
0	8.0	NaN	16.9	21.8	No	No	
1	NaN	NaN	17.2	24.3	No	No	
2	NaN	2.0	21.0	23.2	No	No	
3	NaN	NaN	18.1	26.5	No	No	
4	7.0	8.0	17.8	29.7	No	No	
...	
144729	NaN	NaN	8.8	17.2	No	No	
144730	NaN	NaN	7.0	15.7	No	No	
144731	NaN	3.0	8.9	18.0	No	No	
144732	6.0	7.0	11.7	21.5	No	No	
144733	3.0	NaN	10.6	20.2	No	No	

```
Month Season
```



```

0      December  Summer
1      December  Summer
2      December  Summer
3      December  Summer
4      December  Summer
...
144729      June   Winter
144730      June   Winter
144731      June   Winter
144732      June   Winter
144733      June   Winter

```

[107502 rows x 25 columns]

```
[121]: X_train_enc, y_train, X_test_enc, y_test, preprocessor = preprocess_features(
        train_df,
        test_df,
        numeric_features,
        categorical_features + ["Season"],
        drop_features + ["Month"],
    )
```

```
[122]: X_train_enc.columns
```

```
[122]: Index(['MinTemp', 'MaxTemp', 'Rainfall', 'Evaporation', 'Sunshine',
            'WindGustSpeed', 'WindSpeed9am', 'WindSpeed3pm', 'Humidity9am',
            'Humidity3pm',
            ...,
            'x3_W', 'x3_WNW', 'x3_WSW', 'x4_?', 'x4_No', 'x4_Yes', 'x5_Autumn',
            'x5_Fall', 'x5_Summer', 'x5_Winter'],
            dtype='object', length=123)
```

```
[123]: coeff_df = score_lr_print_coeff(
        preprocessor, train_df, y_train, test_df, y_test, X_train_enc.columns
    )
```

Train score: 0.85

Test score: 0.84

```
[124]: coeff_df.loc[["x5_Fall", "x5_Summer", "x5_Winter", "x5_Autumn"]]
```


```
[124]:
```

	Coef
x5_Fall	0.063903
x5_Summer	-0.225374
x5_Winter	0.105379
x5_Autumn	0.044957


- No improvements in the scores but the coefficients make some sense,
- A negative coefficient for summer and a positive coefficients for winter.

Let's explore Date/Rainfall plots

```
[125]: train_df.plot(x="Date", y="Rainfall")  
plt.ylabel("Rainfall (mm)");
```




```
[126]: monthly_avg_rainfall = train_df.groupby("Month")["Rainfall"].mean()  
plt.plot(monthly_avg_rainfall)  
plt.xticks(np.arange(1, 13).astype(int))  
plt.ylabel("Avg rainfall (mm)")  
plt.xlabel("Month")  
plt.xticks(rotation=90);
```




- It's interesting that June is rainy but May and August are less so.
- But, Australia is a huge country. Perhaps we should drill down to particular locations:

```
[127]: train_df_canberra = train_df.query('Location == "Canberra"')
```

```
[128]: train_df_canberra.plot(x="Date", y="Rainfall")  
plt.ylabel("Rainfall (mm)");
```



```
[129]: plt.plot(train_df_canberra.groupby("Month")["Rainfall"].mean())
plt.xticks(np.arange(1, 13).astype(int))
plt.ylabel("Avg rainfall (mm)")
plt.xlabel("Month")
plt.title("Rainfall in Canberra")
plt.xticks(rotation=90);
```



```
[130]: train_df_canberra.shape
```

```
[130]: (2696, 25)
```

- This looks somewhat cleaner but also pretty surprising - why is December so much higher than January?
- Let's find the location with max rainfall

```
[131]: loc_rainfall = train_df.groupby("Location")["Rainfall"].mean()
loc_rainfall.head()
```

```
[131]: Location
Adelaide      1.526183
Albany        2.292733
Albury        1.978531
AliceSprings  0.899654
BadgerysCreek 2.282515
```


Name: Rainfall, dtype: float64

```
[132]: loc_rainfall.idxmax() # location with max rainfall
```


```
[132]: 'Cairns'
```

```
[133]: train_df_cairns = train_df.query('Location == "Cairns"')
```

```
[134]: train_df_cairns.plot(x="Date", y="Rainfall")  
plt.ylabel("Rainfall (mm)");
```



```
[135]: plt.plot(train_df_cairns.groupby("Month")["Rainfall"].mean())  
plt.xticks(np.arange(1, 13).astype(int))  
plt.ylabel("Avg rainfall (mm)")  
plt.xlabel("Month")  
plt.title("Rainfall in Cairns")  
plt.xticks(rotation=90);
```



```
[136]: train_df_cairns.shape
```

```
[136]: (2310, 25)
```

- This looks much cleaner!

1.7 Lag-based features

- In time series data there is **temporal dependence**;
 - observations close in time tend to be correlated.
- Currently we're using features about today to predict tomorrow's rainfall.
- But, what if tomorrow's rainfall is also related to yesterday's features, or the day before?
 - This is called a **lagged feature**.
- In time series analysis, we'd look at something called an [autocorrelation function](#) (ACF), but we won't go into that here.
- Instead, we can just add those features:

```
[137]: train_df = rain_df.query("Date <= 20150630")
test_df = rain_df.query("Date > 20150630")
```

```
[138]: train_df
```

```
[138]:
```

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	\
0	2008-12-01	Albury	13.4	22.9	0.6	NaN	NaN	
1	2008-12-02	Albury	7.4	25.1	0.0	NaN	NaN	
2	2008-12-03	Albury	12.9	25.7	0.0	NaN	NaN	
3	2008-12-04	Albury	9.2	28.0	0.0	NaN	NaN	
4	2008-12-05	Albury	17.5	32.3	1.0	NaN	NaN	
...	
144729	2015-06-26	Uluru	3.8	18.3	0.0	NaN	NaN	
144730	2015-06-27	Uluru	2.5	17.1	0.0	NaN	NaN	
144731	2015-06-28	Uluru	4.5	19.6	0.0	NaN	NaN	
144732	2015-06-29	Uluru	7.6	22.0	0.0	NaN	NaN	
144733	2015-06-30	Uluru	6.8	21.1	0.0	NaN	NaN	

	WindGustDir	WindGustSpeed	WindDir9am	...	Humidity9am	Humidity3pm	\
0	W	44.0	W	...	71.0	22.0	
1	WNW	44.0	NNW	...	44.0	25.0	
2	WSW	46.0	W	...	38.0	30.0	
3	NE	24.0	SE	...	45.0	16.0	
4	W	41.0	ENE	...	82.0	33.0	
...	
144729	E	39.0	ESE	...	73.0	37.0	
144730	E	41.0	ESE	...	69.0	40.0	
144731	ENE	35.0	ESE	...	69.0	39.0	
144732	ESE	33.0	SE	...	67.0	37.0	
144733	ESE	35.0	ESE	...	81.0	35.0	

	Pressure9am	Pressure3pm	Cloud9am	Cloud3pm	Temp9am	Temp3pm	\
0	1007.7	1007.1	8.0	NaN	16.9	21.8	
1	1010.6	1007.8	NaN	NaN	17.2	24.3	
2	1007.6	1008.7	NaN	2.0	21.0	23.2	
3	1017.6	1012.8	NaN	NaN	18.1	26.5	
4	1010.8	1006.0	7.0	8.0	17.8	29.7	
...	
144729	1031.5	1027.6	NaN	NaN	8.8	17.2	
144730	1029.9	1026.0	NaN	NaN	7.0	15.7	
144731	1028.7	1025.0	NaN	3.0	8.9	18.0	
144732	1027.2	1023.8	6.0	7.0	11.7	21.5	
144733	1028.6	1025.2	3.0	NaN	10.6	20.2	

	RainToday	RainTomorrow
0	No	No
1	No	No

2	No	No
3	No	No
4	No	No
...
144729	No	No
144730	No	No
144731	No	No
144732	No	No
144733	No	No

[107502 rows x 23 columns]

- It looks like the dataframe is already sorted by Location and then by date for each Location.
- We could have done this ourselves with:

```
[139]: # train_df.sort_values(by=["Location", "Date"])
```

But make sure to also sort the targets (i.e. do this before preprocessing).

We can “lag” (or “shift”) a time series in Pandas with the `.shift()` method.

```
[140]: train_df = train_df.assign(Rainfall_lag1=train_df["Rainfall"].shift(1))
```

```
[141]: train_df[["Date", "Location", "Rainfall", "Rainfall_lag1"]].head(10)
```

```
[141]:
```

	Date	Location	Rainfall	Rainfall_lag1
0	2008-12-01	Albury	0.6	NaN
1	2008-12-02	Albury	0.0	0.6
2	2008-12-03	Albury	0.0	0.0
3	2008-12-04	Albury	0.0	0.0
4	2008-12-05	Albury	1.0	0.0
5	2008-12-06	Albury	0.2	1.0
6	2008-12-07	Albury	0.0	0.2
7	2008-12-08	Albury	0.0	0.0
8	2008-12-09	Albury	0.0	0.0
9	2008-12-10	Albury	1.4	0.0

- But we have **multiple time series** here and we need to be more careful with this.
- When we switch from one location to another we do not want to take the value from the previous location.

```
[142]: def create_lag_feature(df, orig_feature, lag):
        """Creates a new df with a new feature that's a lagged version of the
        ↪original, where lag is an int."""
        # note: pandas .shift() kind of does this for you already, but oh well I
        ↪already wrote this code

        new_df = df.copy()
        new_feature_name = "%s_lag%d" % (orig_feature, lag)
```

```

new_df[new_feature_name] = np.nan
for location, df_location in new_df.groupby(
    "Location"
):
    # Each location is its own time series
    new_df.loc[df_location.index[lag:], new_feature_name] = df_location.
    ↪iloc[::-lag][
        orig_feature
    ].values
return new_df

```

```
[143]: train_df = create_lag_feature(train_df, "Rainfall", 1)
```

```
[144]: train_df[["Date", "Location", "Rainfall", "Rainfall_lag1"]][2285:2295]
```

```
[144]:
```

	Date	Location	Rainfall	Rainfall_lag1
2309	2015-06-26	Albury	0.2	1.0
2310	2015-06-27	Albury	0.0	0.2
2311	2015-06-28	Albury	0.2	0.0
2312	2015-06-29	Albury	0.0	0.2
2313	2015-06-30	Albury	0.0	0.0
3040	2009-01-01	BadgerysCreek	0.0	NaN
3041	2009-01-02	BadgerysCreek	0.0	0.0
3042	2009-01-03	BadgerysCreek	0.0	0.0
3043	2009-01-04	BadgerysCreek	0.0	0.0
3044	2009-01-05	BadgerysCreek	0.0	0.0

Now it looks good!

- Question: is it OK to do this to the test set? Discuss.
- It's fine if you would have this information available in deployment.
- If we're just forecasting the next day, we should.
- Let's include it for now.

```
[145]: rain_df_modified = create_lag_feature(rain_df, "Rainfall", 1)
train_df = rain_df_modified.query("Date <= 20150630")
test_df = rain_df_modified.query("Date > 20150630")
```

```
[146]: rain_df_modified
```

```
[146]:
```

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	\
0	2008-12-01	Albury	13.4	22.9	0.6	NaN	NaN	
1	2008-12-02	Albury	7.4	25.1	0.0	NaN	NaN	
2	2008-12-03	Albury	12.9	25.7	0.0	NaN	NaN	
3	2008-12-04	Albury	9.2	28.0	0.0	NaN	NaN	
4	2008-12-05	Albury	17.5	32.3	1.0	NaN	NaN	
...	
145454	2017-06-20	Uluru	3.5	21.8	0.0	NaN	NaN	
145455	2017-06-21	Uluru	2.8	23.4	0.0	NaN	NaN	

145456	2017-06-22	Uluru	3.6	25.3	0.0	NaN	NaN
145457	2017-06-23	Uluru	5.4	26.9	0.0	NaN	NaN
145458	2017-06-24	Uluru	7.8	27.0	0.0	NaN	NaN

	WindGustDir	WindGustSpeed	WindDir9am	...	Humidity3pm	Pressure9am	\
0	W	44.0	W	...	22.0	1007.7	
1	WNW	44.0	NNW	...	25.0	1010.6	
2	WSW	46.0	W	...	30.0	1007.6	
3	NE	24.0	SE	...	16.0	1017.6	
4	W	41.0	ENE	...	33.0	1010.8	
...	
145454	E	31.0	ESE	...	27.0	1024.7	
145455	E	31.0	SE	...	24.0	1024.6	
145456	NNW	22.0	SE	...	21.0	1023.5	
145457	N	37.0	SE	...	24.0	1021.0	
145458	SE	28.0	SSE	...	24.0	1019.4	

	Pressure3pm	Cloud9am	Cloud3pm	Temp9am	Temp3pm	RainToday	\
0	1007.1	8.0	NaN	16.9	21.8	No	
1	1007.8	NaN	NaN	17.2	24.3	No	
2	1008.7	NaN	2.0	21.0	23.2	No	
3	1012.8	NaN	NaN	18.1	26.5	No	
4	1006.0	7.0	8.0	17.8	29.7	No	
...		
145454	1021.2	NaN	NaN	9.4	20.9	No	
145455	1020.3	NaN	NaN	10.1	22.4	No	
145456	1019.1	NaN	NaN	10.9	24.5	No	
145457	1016.8	NaN	NaN	12.5	26.1	No	
145458	1016.5	3.0	2.0	15.1	26.0	No	

	RainTomorrow	Rainfall_lag1
0	No	NaN
1	No	0.6
2	No	0.0
3	No	0.0
4	No	0.0
...
145454	No	0.0
145455	No	0.0
145456	No	0.0
145457	No	0.0
145458	No	0.0

[142193 rows x 24 columns]

```
[147]: X_train_enc, y_train, X_test_enc, y_test, preprocessor = preprocess_features(
        train_df,
```

```

test_df,
numeric_features + ["Rainfall_lag1"],
categorical_features,
drop_features,
)

```

```

[148]: lr_coef = score_lr_print_coef(
        preprocessor, train_df, y_train, test_df, y_test, X_train_enc.columns
    )
    # lr_coef

```

Train score: 0.85

Test score: 0.84

```

[149]: lr_coef.loc[["Rainfall", "Rainfall_lag1"]]

```

```

[149]:
          Coef
Rainfall    0.081068
Rainfall_lag1 0.008305

```

- Rainfall from today has a positive coefficient.
- Rainfall from yesterday has a positive but a smaller coefficient.
- If we **didn't have rainfall from today** feature, rainfall from **yesterday** feature would have received a **bigger coefficient**.
- We could also create a lagged version of the target.
- In fact, this dataset already has that built in! **RainToday** is the lagged version of the target **RainTomorrow**.
- We could also create lagged version of other features, or more lags

```

[150]: rain_df_modified = create_lag_feature(rain_df, "Rainfall", 1)
rain_df_modified = create_lag_feature(rain_df_modified, "Rainfall", 2)
rain_df_modified = create_lag_feature(rain_df_modified, "Rainfall", 3)
rain_df_modified = create_lag_feature(rain_df_modified, "Humidity3pm", 1)

```

```

[151]: rain_df_modified[
        [
            "Date",
            "Location",
            "Rainfall",
            "Rainfall_lag1",
            "Rainfall_lag2",
            "Rainfall_lag3",
            "Humidity3pm",
            "Humidity3pm_lag1",
        ]
    ].head(10)

```

```
[151]:
```

	Date	Location	Rainfall	Rainfall_lag1	Rainfall_lag2	Rainfall_lag3	\
0	2008-12-01	Albury	0.6	NaN	NaN	NaN	
1	2008-12-02	Albury	0.0	0.6	NaN	NaN	
2	2008-12-03	Albury	0.0	0.0	0.6	NaN	
3	2008-12-04	Albury	0.0	0.0	0.0	0.6	
4	2008-12-05	Albury	1.0	0.0	0.0	0.0	
5	2008-12-06	Albury	0.2	1.0	0.0	0.0	
6	2008-12-07	Albury	0.0	0.2	1.0	0.0	
7	2008-12-08	Albury	0.0	0.0	0.2	1.0	
8	2008-12-09	Albury	0.0	0.0	0.0	0.2	
9	2008-12-10	Albury	1.4	0.0	0.0	0.0	

	Humidity3pm	Humidity3pm_lag1
0	22.0	NaN
1	25.0	22.0
2	30.0	25.0
3	16.0	30.0
4	33.0	16.0
5	23.0	33.0
6	19.0	23.0
7	19.0	19.0
8	9.0	19.0
9	27.0	9.0

Note the pattern of NaN values.

```
[152]: train_df = rain_df_modified.query("Date <= 20150630")
test_df = rain_df_modified.query("Date > 20150630")
```

```
[153]: X_train_enc, y_train, X_test_enc, y_test, preprocessor = preprocess_features(
    train_df,
    test_df,
    numeric_features
    + ["Rainfall_lag1", "Rainfall_lag2", "Rainfall_lag3", "Humidity3pm_lag1"],
    categorical_features,
    drop_features,
)
```

```
[154]: lr_coef = score_lr_print_coef(
    preprocessor, train_df, y_train, test_df, y_test, X_train_enc.columns
)
```

Train score: 0.85

Test score: 0.85

```
[155]: lr_coef.loc[
    [
        "Rainfall",
```

```

    "Rainfall_lag1",
    "Rainfall_lag2",
    "Rainfall_lag3",
    "Humidity3pm",
    "Humidity3pm_lag1",
]
]

```

```

[155]:
           Coef
Rainfall      0.108510
Rainfall_lag1 0.023072
Rainfall_lag2 0.018270
Rainfall_lag3 0.017805
Humidity3pm    1.278646
Humidity3pm_lag1 -0.267520

```

Note the **pattern in the magnitude** of the coefficients.

1.8 Forecasting further into the future

- Let's say we want to predict 7 days into the future instead of one day.
- There are a few main approaches here:
 1. Train a separate model for **each number of days**.
 - E.g. one model that predicts RainTomorrow, another model that predicts RainIn2Days, etc. We can build these datasets.
 2. Use a **multi-output model** that jointly predicts RainTomorrow, RainIn2Days, etc. However, multi-output models are outside the scope of CPSC 330.
 3. Use **one model** and **sequentially predict** using a for loop.
 - However, this **requires predicting *all* features** into a model so may not be that useful here.
- To briefly dig into approach 3, this is easier to understand for a univariate (one feature) time series.
- To dig into this we'll look at the [Retail Sales of Clothing and Clothing Accessory Stores dataset](#) made available by the Federal Reserve Bank of St. Louis.

```

[156]: retail_df = pd.read_csv("data/MRTSSM448USN.csv", parse_dates=["DATE"])
       retail_df.columns = ["date", "sales"]

```

```

[157]: retail_df.head()

```

```

[157]:
   date  sales
0 1992-01-01  6938
1 1992-02-01  7524
2 1992-03-01  8475
3 1992-04-01  9401
4 1992-05-01  9558

```

```
[158]: retail_df["date"].min()
```


```
[158]: Timestamp('1992-01-01 00:00:00')
```

```
[159]: retail_df["date"].max()
```

```
[159]: Timestamp('2022-04-01 00:00:00')
```

```
[160]: retail_df_train = retail_df.query("date <= 20170101")
retail_df_test = retail_df.query("date > 20170101")
```

```
[161]: retail_df_train.plot(x="date", y="sales", figsize=(15, 5));
```



We can create a dataset using purely lag features.

```
[162]: def lag_df(df, lag, cols):
        return df.assign(
            **{f"{col}-{n}": df[col].shift(n) for n in range(1, lag + 1) for col in cols}
        )
```

```
[163]: retail_lag_5 = lag_df(retail_df, 5, ["sales"])
retail_train_5 = retail_lag_5.query("date <= 20170101")
retail_test_5 = retail_lag_5.query("date > 20170101")
retail_train_5
```

```
[163]:
```

	date	sales	sales-1	sales-2	sales-3	sales-4	sales-5
0	1992-01-01	6938	NaN	NaN	NaN	NaN	NaN
1	1992-02-01	7524	6938.0	NaN	NaN	NaN	NaN
2	1992-03-01	8475	7524.0	6938.0	NaN	NaN	NaN
3	1992-04-01	9401	8475.0	7524.0	6938.0	NaN	NaN
4	1992-05-01	9558	9401.0	8475.0	7524.0	6938.0	NaN
..

296	2016-09-01	19928	22505.0	20782.0	20274.0	21774.0	20514.0
297	2016-10-01	20650	19928.0	22505.0	20782.0	20274.0	21774.0
298	2016-11-01	23826	20650.0	19928.0	22505.0	20782.0	20274.0
299	2016-12-01	34847	23826.0	20650.0	19928.0	22505.0	20782.0
300	2017-01-01	15921	34847.0	23826.0	20650.0	19928.0	22505.0

[301 rows x 7 columns]

- Now, if we drop the “date” column we have
 - a target (“sales”) and
 - 5 features (the previous 5 days of sales).
- We need to impute/drop the missing values and then we can fit a model to this. I will just drop for convenience:

```
[164]: retail_train_5 = retail_train_5.drop(columns=["date"])
       retail_train_5
```

```
[164]:
```

	sales	sales-1	sales-2	sales-3	sales-4	sales-5
5	9182	9558.0	9401.0	8475.0	7524.0	6938.0
6	9103	9182.0	9558.0	9401.0	8475.0	7524.0
7	10513	9103.0	9182.0	9558.0	9401.0	8475.0
8	9573	10513.0	9103.0	9182.0	9558.0	9401.0
9	10254	9573.0	10513.0	9103.0	9182.0	9558.0
..
296	19928	22505.0	20782.0	20274.0	21774.0	20514.0
297	20650	19928.0	22505.0	20782.0	20274.0	21774.0
298	23826	20650.0	19928.0	22505.0	20782.0	20274.0
299	34847	23826.0	20650.0	19928.0	22505.0	20782.0
300	15921	34847.0	23826.0	20650.0	19928.0	22505.0

[296 rows x 6 columns]

```
[165]: retail_train_5_X = retail_train_5.drop(columns=["sales"])
       retail_train_5_y = retail_train_5["sales"]
```

```
[166]: from sklearn.ensemble import RandomForestRegressor
```

```
[167]: retail_model = RandomForestRegressor()
       retail_model.fit(retail_train_5_X, retail_train_5_y);
```

Given this, we can now predict the sales

```
[168]: preds = retail_model.predict(retail_test_5.drop(columns=["date", "sales"]))
       preds
```

```
[168]: array([18616.48, 20415.87, 21694.89, 22467.23, 21114.88, 22269.33,
              22125.53, 22845.07, 20989.31, 22820.98, 30887.77, 15881.83,
              18627.31, 21130.62, 22240.74, 21868.71, 28765.65, 21034.67,
```



```
22093.24, 29779.64, 21708.63, 21985.46, 21486.28, 15932. ,
18772.01, 20404.79, 22391.84, 22212.39, 28360.13, 22081.68,
21925.26, 30677.25, 21577.89, 21948.33, 27823.64, 15941.43,
19031.6 , 21439.78, 13157.29, 11188.87, 10875.16, 14550.56,
11912.13, 9623.18, 18518.62, 20691.91, 23083.33, 14687.02,
17664.73, 18970.94, 31352.74, 29340.92, 15990.5 , 28299.48,
26286.69, 28365.17, 23821.67, 29063.58, 16112.38, 16522.07,
20382.02, 22711.27, 19152.86])
```

```
[169]: retail_test_5_preds = retail_test_5.assign(predicted_sales=preds)
       retail_test_5_preds.head()
```

```
[169]:      date  sales  sales-1  sales-2  sales-3  sales-4  sales-5  \
301 2017-02-01  18036  15921.0  34847.0  23826.0  20650.0  19928.0
302 2017-03-01  21348  18036.0  15921.0  34847.0  23826.0  20650.0
303 2017-04-01  21154  21348.0  18036.0  15921.0  34847.0  23826.0
304 2017-05-01  21954  21154.0  21348.0  18036.0  15921.0  34847.0
305 2017-06-01  20623  21954.0  21154.0  21348.0  18036.0  15921.0


      predicted_sales
301          18616.48
302          20415.87
303          21694.89
304          22467.23
305          21114.88
```

- Ok, that is fine, but what if we want to predict 7 days in the future?
 - Well, we would not have access to our features!! We don't yet know the previous day's sales, or 2 days prior!
 - So we can use "Approach 3" mentioned earlier:
 - predict these values and then pretend they are true!
 - For simplicity, say today is Monday
1. **Predict** Tuesday's sales
 2. Then, to predict for Wednesday, we need to know Tuesday's sales. Use our *prediction* for Tuesday as the truth.
 3. Then, to predict for Thursday, we need to know Tue and Wed sales. Use our predictions.
 4. Etc etc.

1.9 Trends

- There are some important concepts in time series that rely on having a continuous target (like we do in the retail sales example above).
- Part of that is the idea of seasonality and trends.
- These are mostly taken care of by our feature engineering of the data variable, but there's something important left to discuss.

```
[170]: retail_df_train.plot(x="date", y="sales", figsize=(15, 5));
```



- It looks like there's a **trend** here - the sales are going up over time.

Let's say we encoded the date as a feature in days like this:

```
[171]: retail_train_5_date = retail_lag_5.query("date <= 20170101")
first_day_retail = retail_train_5_date["date"].min()

retail_train_5_date = retail_train_5_date.assign(
    Days_since=retail_train_5_date["date"].apply(lambda x: (x -
↪first_day_retail).days)
)
retail_train_5_date.head(10)
```

```
[171]:
```

	date	sales	sales-1	sales-2	sales-3	sales-4	sales-5	Days_since
0	1992-01-01	6938	NaN	NaN	NaN	NaN	NaN	0
1	1992-02-01	7524	6938.0	NaN	NaN	NaN	NaN	31
2	1992-03-01	8475	7524.0	6938.0	NaN	NaN	NaN	60
3	1992-04-01	9401	8475.0	7524.0	6938.0	NaN	NaN	91
4	1992-05-01	9558	9401.0	8475.0	7524.0	6938.0	NaN	121
5	1992-06-01	9182	9558.0	9401.0	8475.0	7524.0	6938.0	152
6	1992-07-01	9103	9182.0	9558.0	9401.0	8475.0	7524.0	182
7	1992-08-01	10513	9103.0	9182.0	9558.0	9401.0	8475.0	213
8	1992-09-01	9573	10513.0	9103.0	9182.0	9558.0	9401.0	244
9	1992-10-01	10254	9573.0	10513.0	9103.0	9182.0	9558.0	274

- Now, let's say we use all these features (the lagged version of the target and also Days_since.
- If we use **linear regression** we'll learn a coefficient for Days_since.
 - If that coefficient is positive, it predicts unlimited growth forever. That may not be what you want? It depends.
- If we use a **random forest**, we'll just be doing splits from the training set, e.g. "if Days_since > 9100 then do this".
 - There will be no splits for later time points because there is no training data there.
 - Thus **tree-based models cannot model trends**.
 - This is really important to know!!

- Often, we **model the trend separately** and
 - use the **random forest** to model a **de-trended** time series.

1.10 What did we not cover?

- A huge amount!

1.10.1 Traditional time series approaches

- Time series analysis is a huge field of its own
- Traditional approaches include the [ARIMA model](#) and its various components/extensions.
- In Python, the [statsmodels](#) package is the place to go for this sort of thing.
 - For example, [statsmodels.tsa.arima_model.ARIMA](#).
- These **approaches can forecast**, and
 - they are also very good for understanding the **temporal relationships** in your data.
- We will take a different route in this course, and **stick to our supervised learning** tools.

1.10.2 Deep learning

- Recently, deep learning has been very successful too.
- In particular, [recurrent neural networks](#) (RNNs).
 - These are not covered in CPSC 340, but I believe they are in 540 (soon to be renamed 440).
 - [LSTMs](#) especially have shown a lot of promise in this type of task.
 - [Here](#) is a blog post about LSTMs.

1.10.3 Types of problems involving time series

- A **single label** associated with an entire time series.
 - We had that with images earlier on, you could have the same for a time series.
 - E.g., for fraud detection, labelling each transaction as fraud/normal vs. labelling a person as bad/good based on their entire history.
 - There are various approaches that can be used for this type of problem, including CNNs (Lecture 14), LSTMs, and non deep learning methods.
- **Inference** problems.
 - What are the **patterns in this time series**?
 - How **many lags** are associated with the current value?
- Etc.

Unequally spaced time points

- We assumed we have a **measurement each day**.
- For example, when creating lag features we used consecutive rows in the DataFrame.
- But, in fact **some days were missing** in this dataset.
- More generally, what if the measurements are at arbitrary times, not equally spaced?
 - Some of our approaches would **still work**, like encoding the month / looking at seasonality.
 - Some of our approaches would **not make sense**, like the lags.
 - Perhaps the measurements could be binned into **equally spaced bins**, or something.
 - This is more of a hassle.

Other software package

- One good one to know about is [Prophet](#).

Feature engineering

- Often, a useful approach is to just *engineer your own features*.
 - E.g., max expenditure, min expenditure, max-min, avg time gap between transactions, variance of time gap between transactions, etc etc.
 - We could do that here as well, or in any problem.